

# Outline

## Beyond ImageNet

1. Fully Convolutional Networks (FCNs)
2. Segmentation
- 3. Transfer learning and Domain adaptation**

# Transfer from ImageNet (source)

## Transfer as generic features

Brut Deep features (learned from ImageNet)

(== a learned embedding from Image to vector representation)

Retrieval



## Transfer learning (from source to target)

Frozen features + SVM => solution to small datasets

Frozen features + Deep

Fine tuning not easy in that case (small datasets)

# Transfer from source(=ImageNet task) to target task

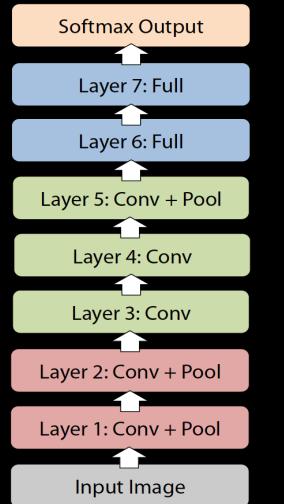
**Source:** ImageNet (dataset + 100 classes) => AlexNet trained

**Target:** new dataset Cal-101 and new classification task with 101 classes => Chopped

AlexNet (layer i) + SVM trained on

## Architecture of Krizhevsky et al.

- 8 layers total
- Trained on Imagenet dataset [Deng et al. CVPR'09]
- 18.2% top-5 error
- Our reimplementation:  
18.1% top-5 error



## Tapping off Features at each Layer

Plug features from each layer into linear SVM or soft-max

	Cal-101 (30/class)	Cal-256 (60/class)
SVM (1)	$44.8 \pm 0.7$	$24.6 \pm 0.4$
SVM (2)	$66.2 \pm 0.5$	$39.6 \pm 0.3$
SVM (3)	$72.3 \pm 0.4$	$46.0 \pm 0.3$
SVM (4)	$76.6 \pm 0.4$	$51.3 \pm 0.1$
SVM (5)	<b><math>86.2 \pm 0.8</math></b>	$65.6 \pm 0.3$
SVM (7)	<b><math>85.5 \pm 0.4</math></b>	<b><math>71.7 \pm 0.2</math></b>
Softmax (5)	$82.9 \pm 0.4$	$65.7 \pm 0.5$
Softmax (7)	<b><math>85.4 \pm 0.4</math></b>	<b><math>72.6 \pm 0.1</math></b>

=> Results better than SoA CV methods on Cal-101!

# Transfer: fine-tuning of a deep model on target task

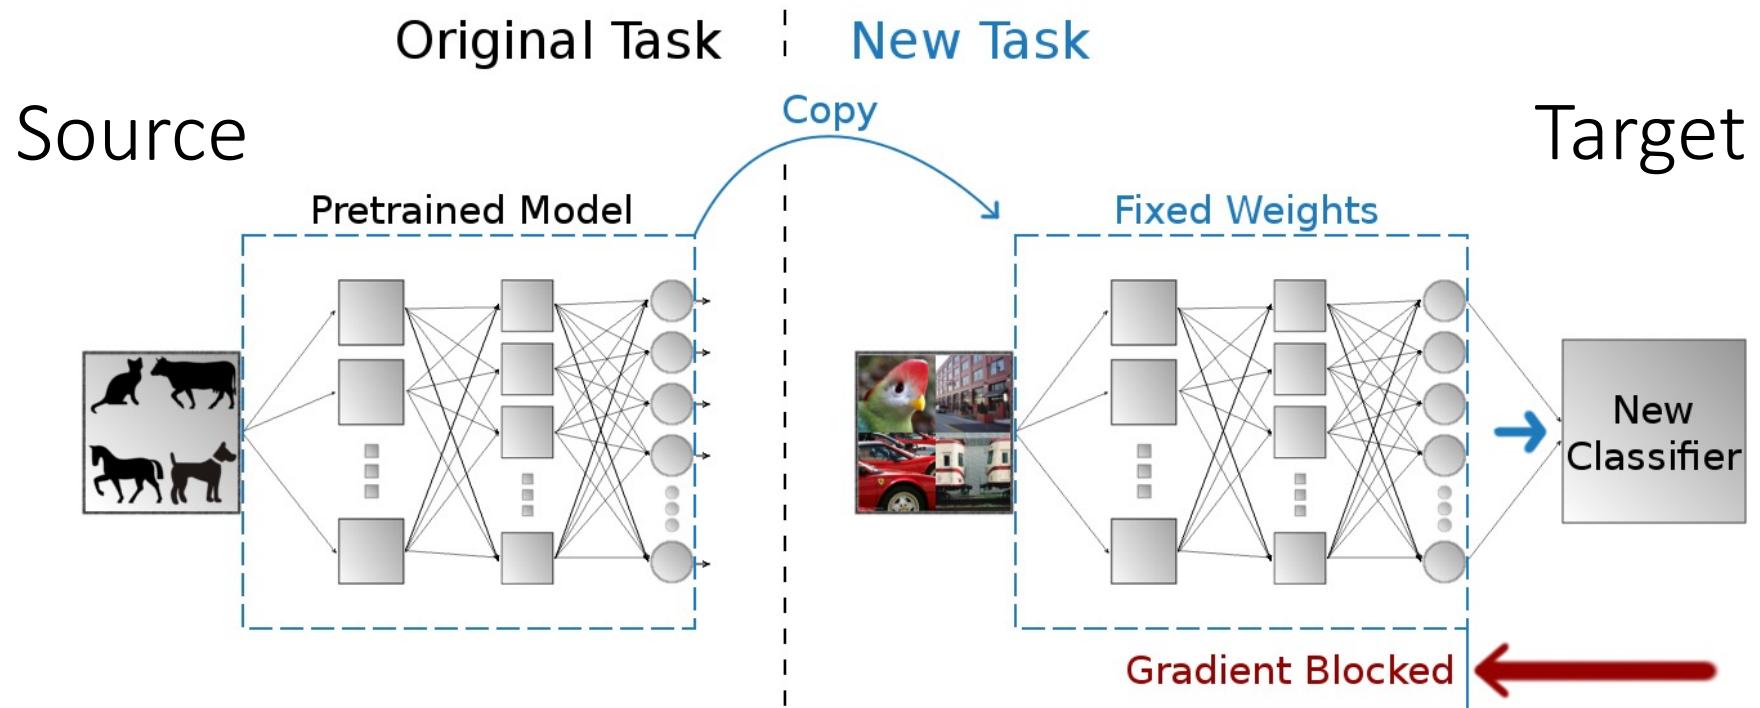
Train a deep (AlexNet) on source (ImageNet)

Keep the deep params. for target and complete with a small deep on top (fully trained on target task)

Fine-tune the whole model on target data

Challenge: only limited target data, careful about overfitting

Solution: Freeze the gradient's update for AlexNet part



# Transfer: fine-tuning of a deep model on target task

Train a deep (AlexNet) on source (ImageNet)

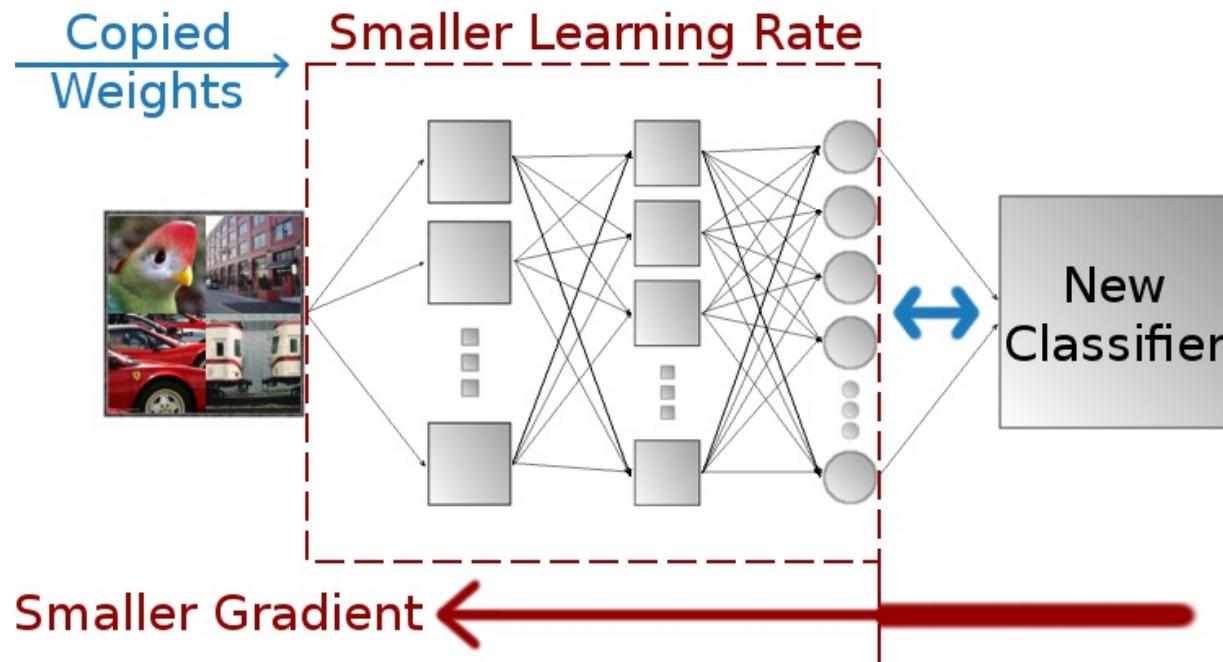
Keep the deep params. for target and complete with a small deep on top (fully trained on target task)

Fine-tune the whole model on target data

Challenge: only limited target data, careful about overfitting

Solution: Freeze the gradient's update for AlexNet part

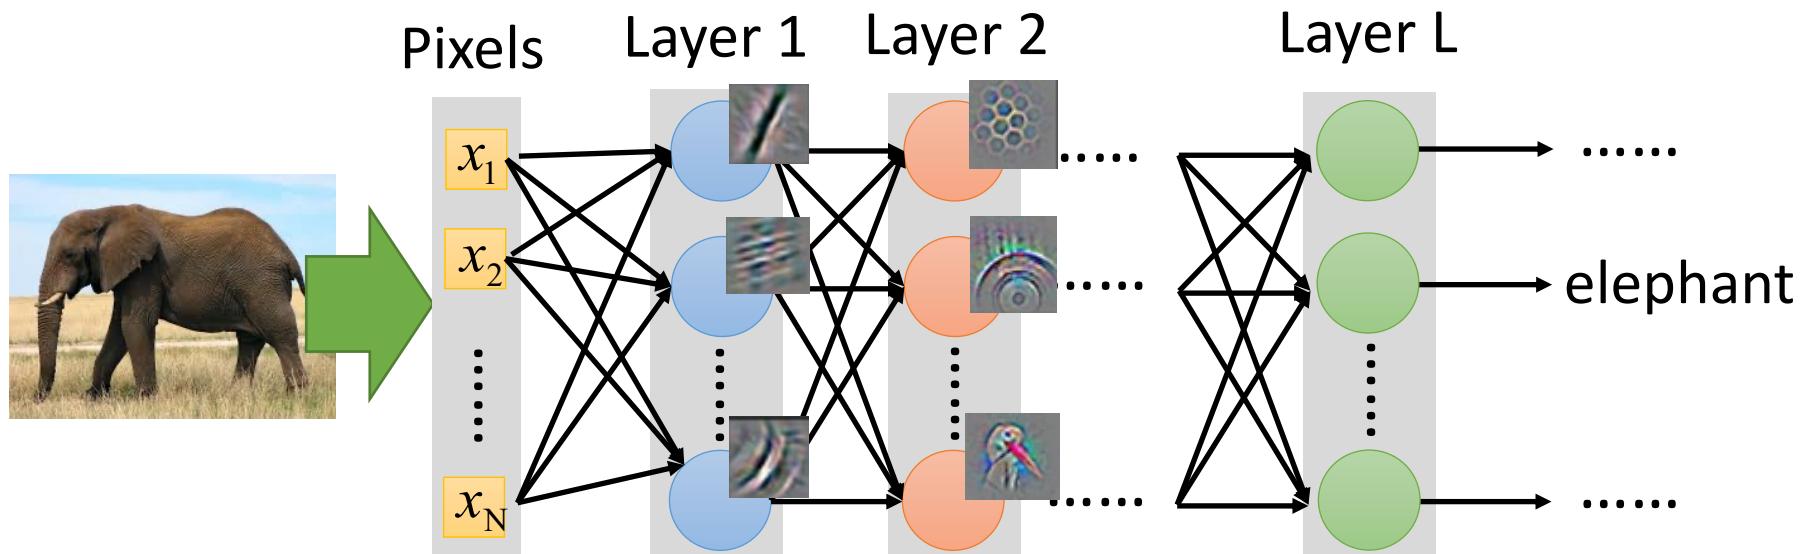
Other solution: use smaller gradient's update for AlexNet part



# Transfer: which parts of the deep?

Which layer(s) can be transferred (copied)?

- Speech: usually copy the last few layers
- Image: usually copy the first few layers



# Transfer: which supervision?

- Task description
  - Source data:  $(x^s, y^s)$   A large amount
  - Target data:  $(x^t, y^t)$   (Very) little

Rq: Few/One-shot learning: only a few/one examples in target domain

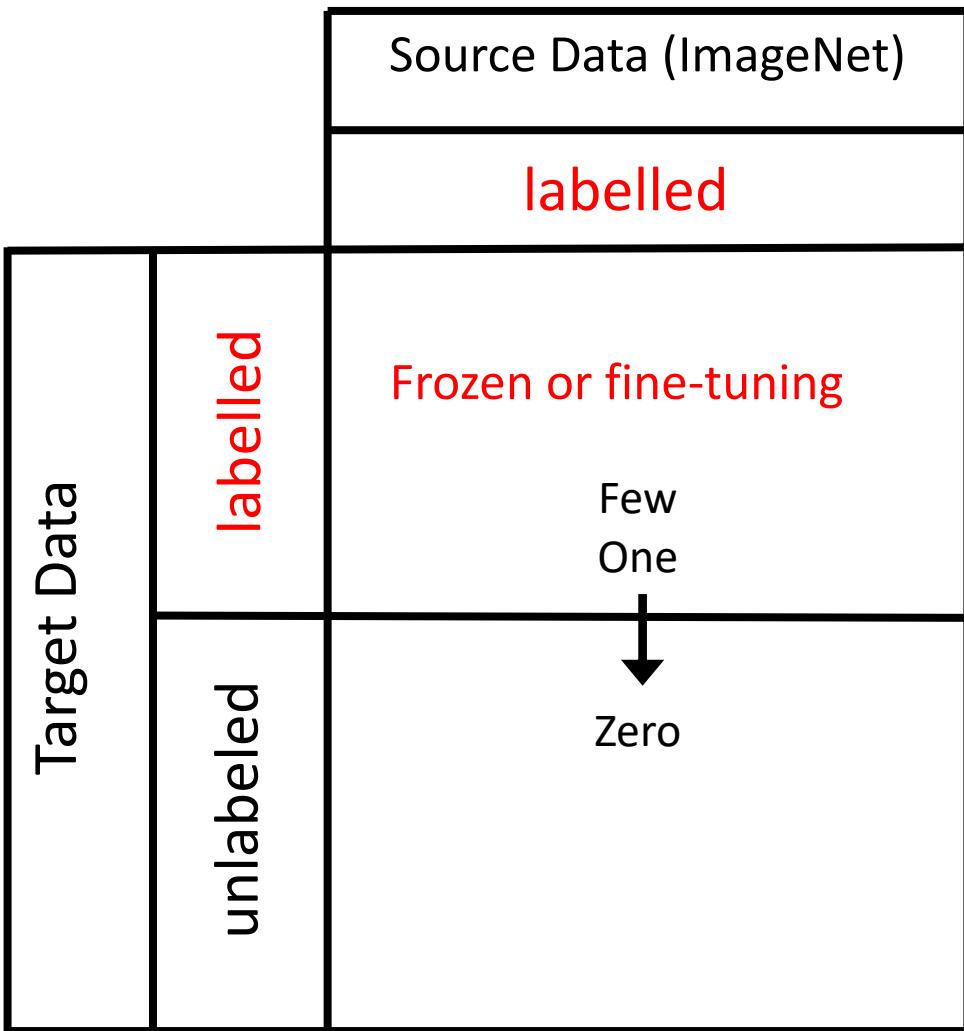
Many different contexts:

In vision: from large dataset ([ImageNet](#)) to small datasets ([VOC2007](#))

In speech: (supervised) speaker adaption

- Source data: audio data and transcriptions from many speakers
- Target data: audio data and its transcriptions of specific user

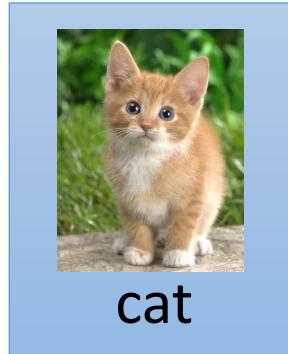
# More on transfer framework



Main purposes:  
Similar visual domain?  
Same tasks (ie class)?

# Similar domain: ImageNet task => Dog/Cat task

Target:  
Dog/Cat  
Classifier



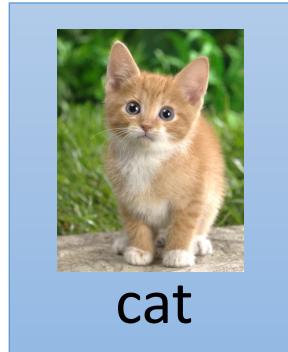
Data *not directly related to* the task considered



ImageNet: Similar domain,  
different task (1000 classes but NOT Dog and Cat classes)

# General Framework for Transfer Learning

Target:  
Dog/Cat  
Classifier



Data *not directly related to* the task considered



elephant

tiger



dog

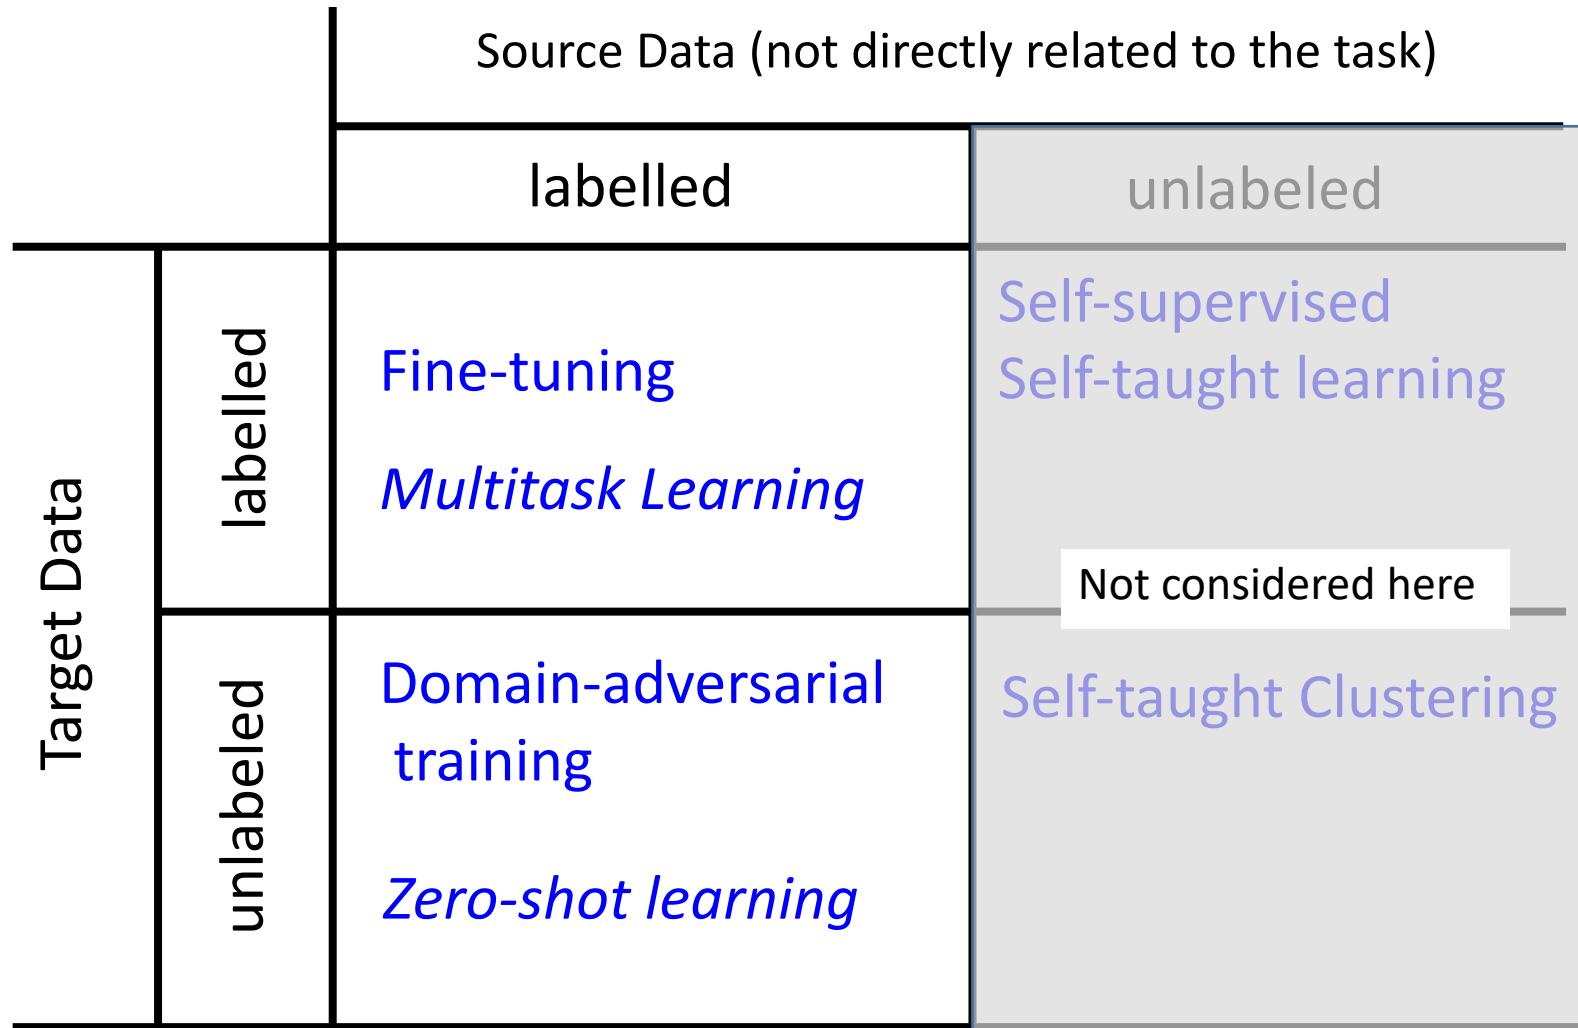


cat

Similar domain, completely  
different tasks

Different domains, same task

# General Framework for Transfer Learning

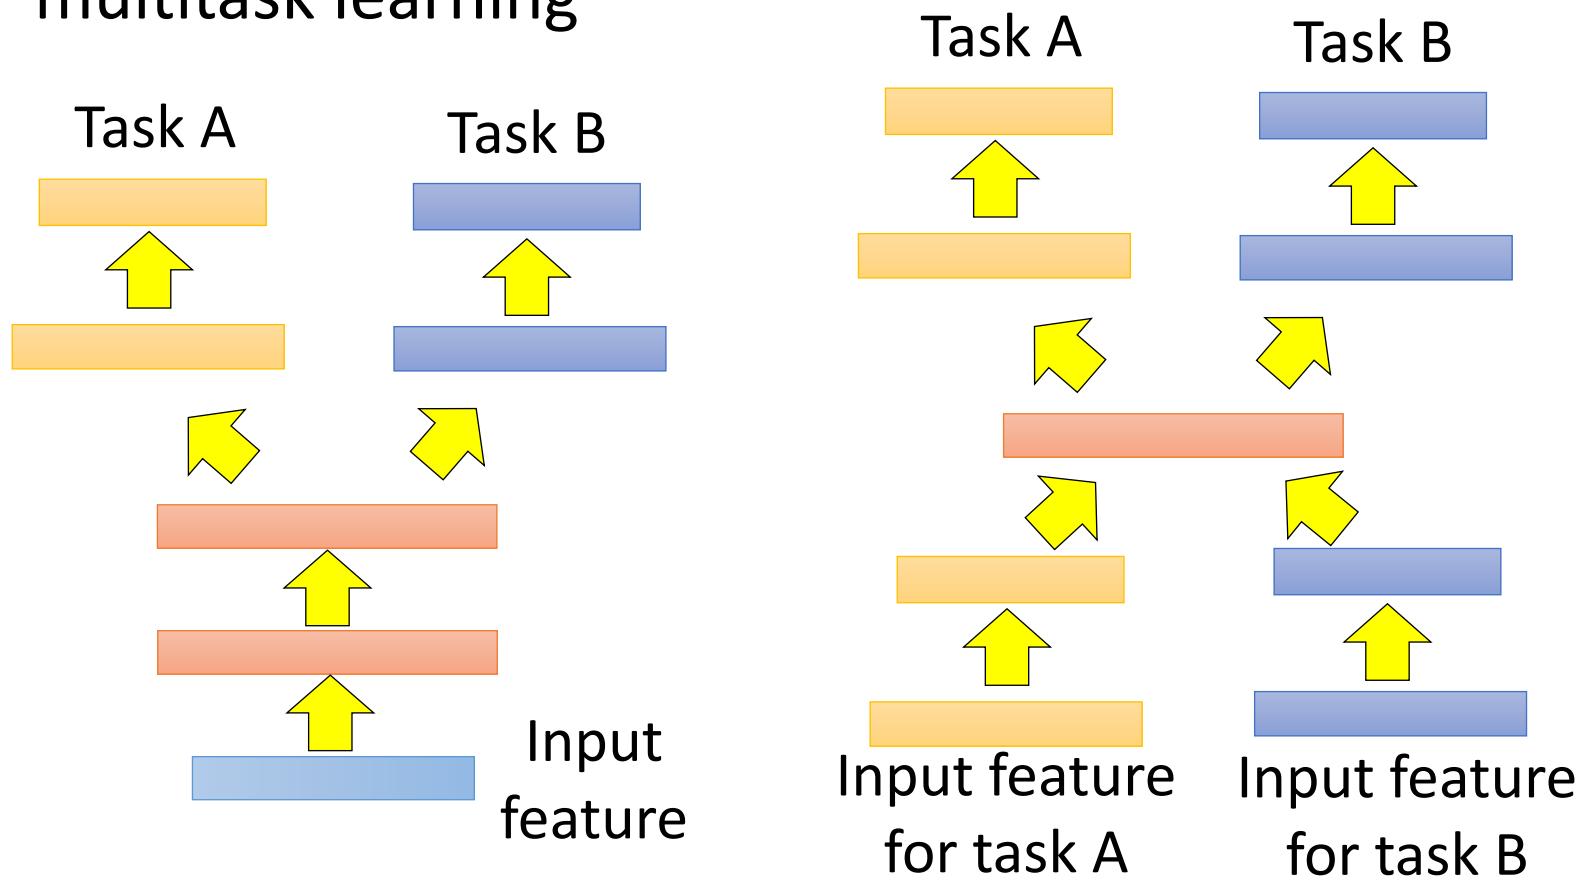


# General Framework for Transfer Learning

		Source Data (not directly related to the task)	
		labelled	unlabeled
Target Data	labelled	Fine-tuning <b><i>Multitask Learning</i></b>	Not considered here
	unlabeled		Not considered here

# Multitask Learning

- The multi-layer structure makes NN suitable for multitask learning



# Transfer Learning - Overview

		Source Data (not directly related to the task)	
		labelled	unlabeled
Target Data	labelled	Fine-tuning <i>Multitask Learning</i>	Not considered here
	unlabeled	<b>Domain adaptation-adversarial training</b>	Not considered here

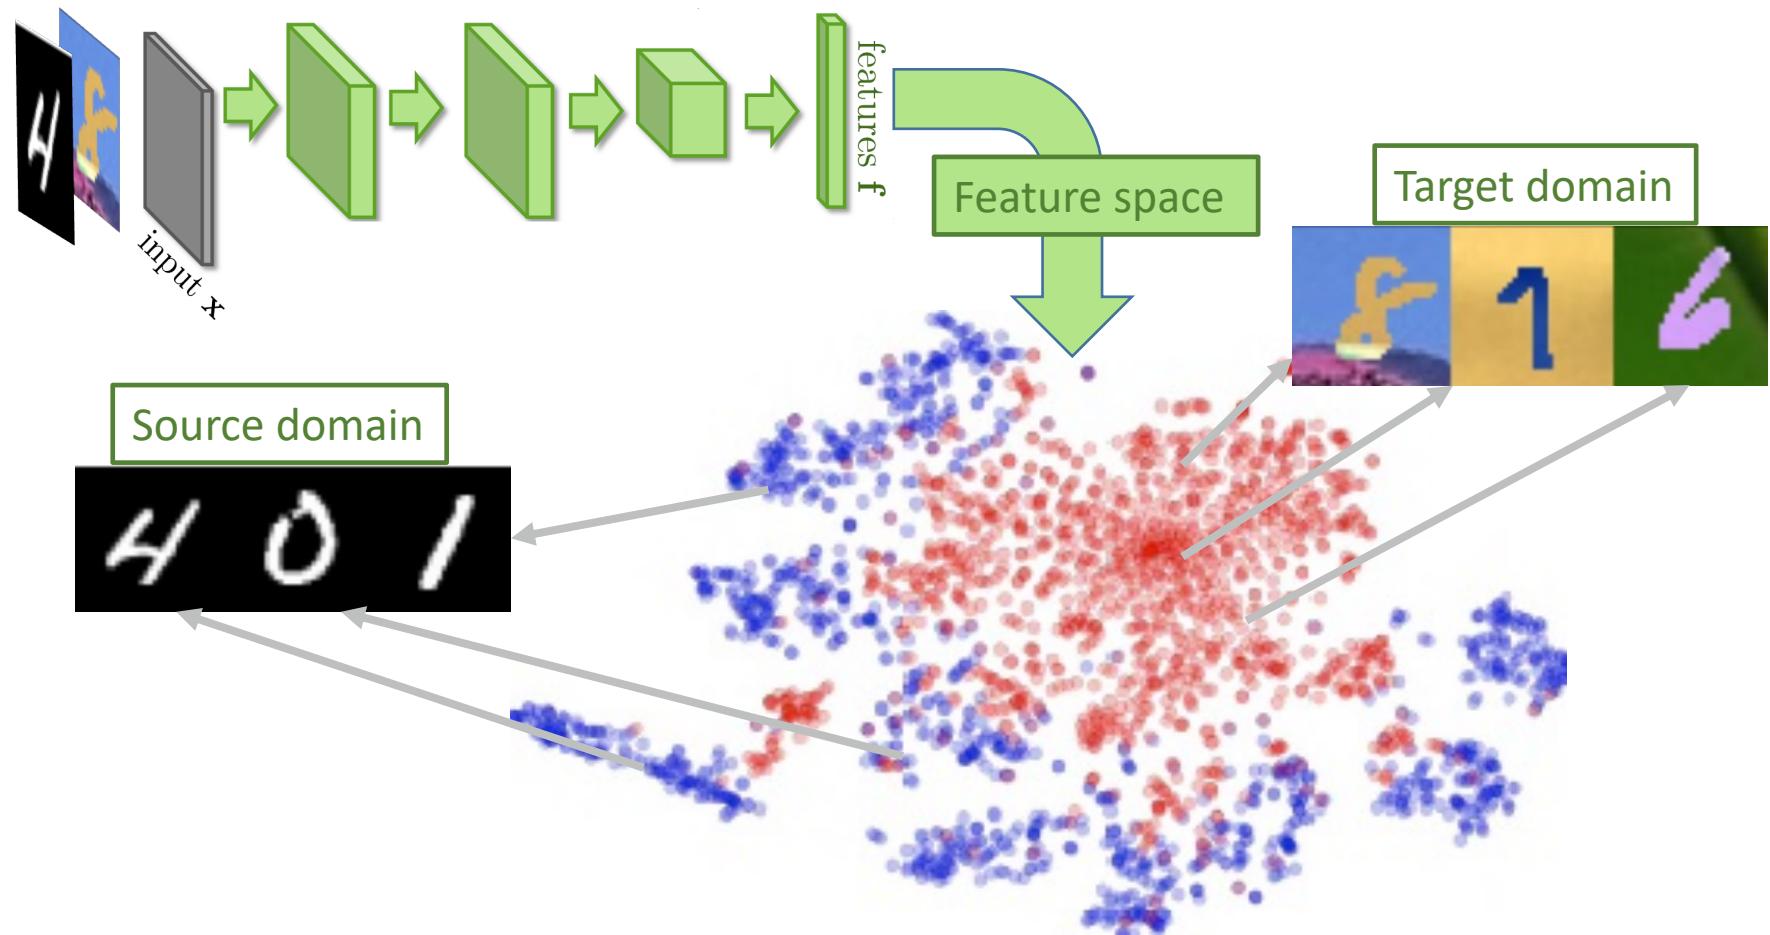
# Unsupervised Domain Adaptation (UDA)

Source data:  $(x^s, y^s)$       Target data:  $(x^t)$        $\xrightarrow{\quad}$  Training data } Same task,  
domain mismatch



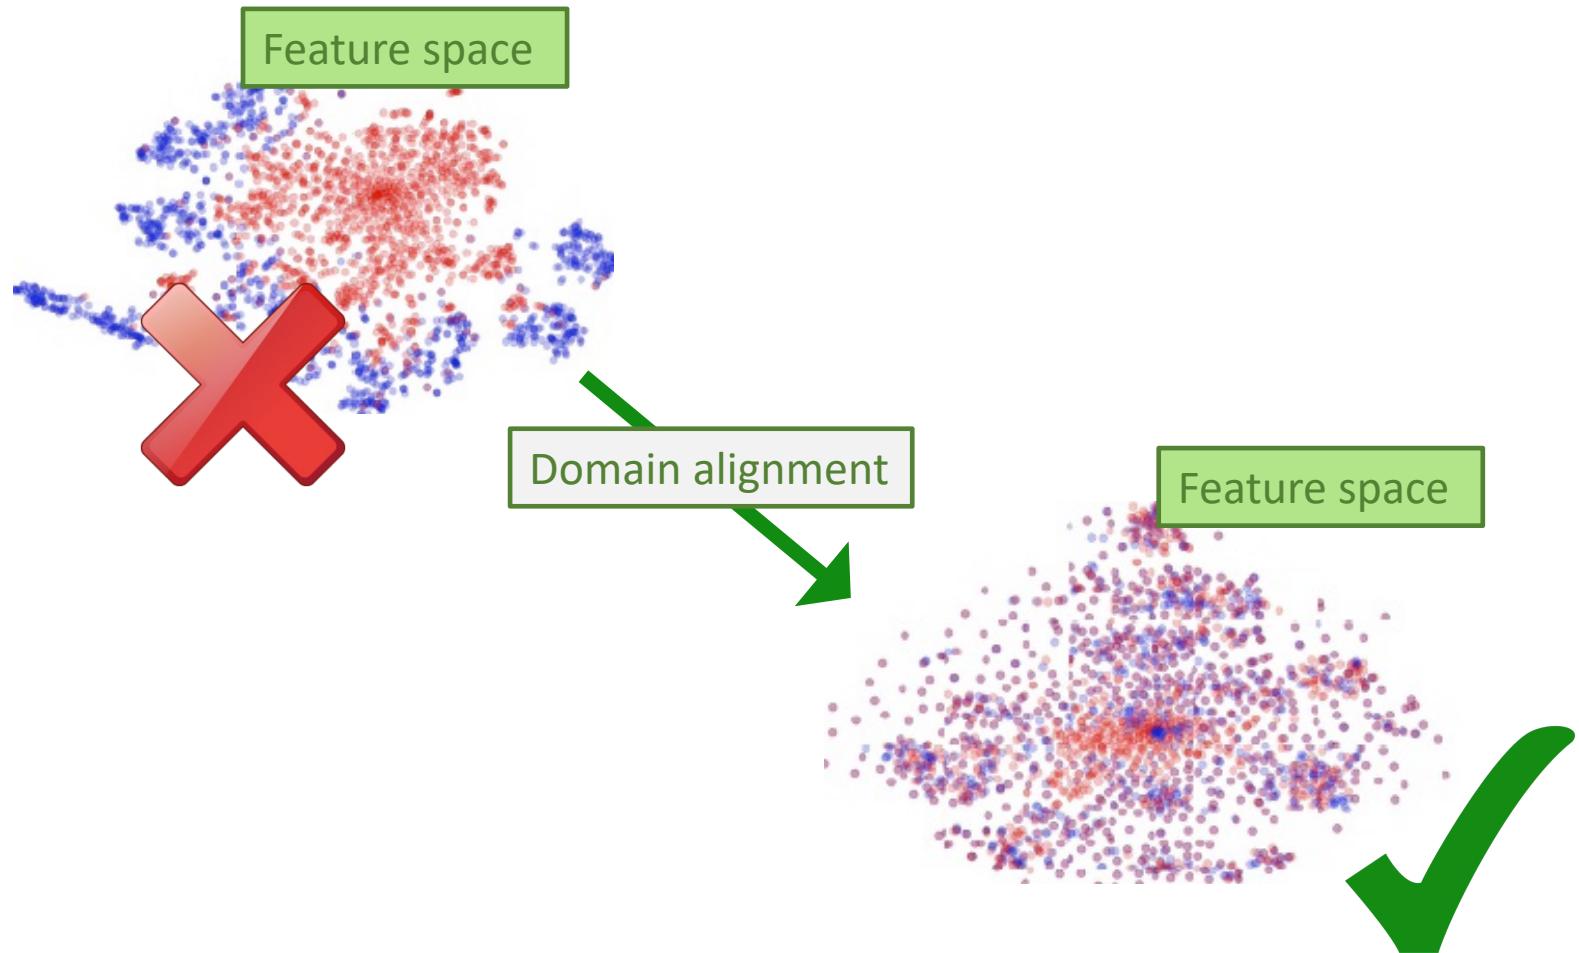
Final test on target domain!

# Unsupervised Domain adaptation (UDA): objectives



Main principle: diminish the **domain** shift in the learned features, encourage domain confusion

# UDA strategy: align both domains



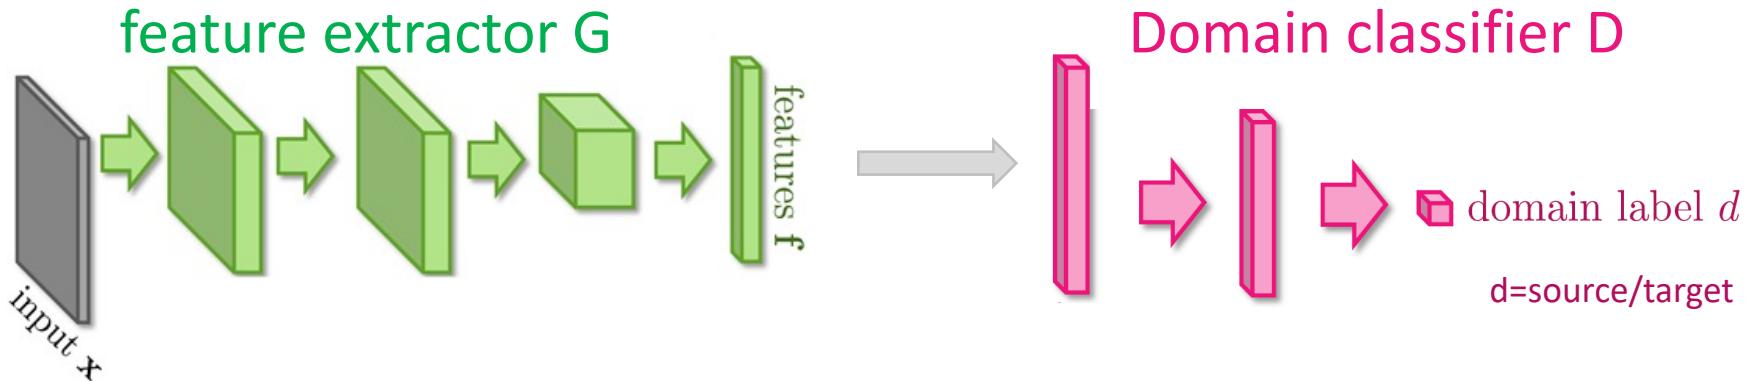
# UDA strategy: 1/ domain-adversarial training

Add to the feature generator (G) a domain classifier (discriminant D) for which labels are available!

Learn G and D:

G tries to align domains

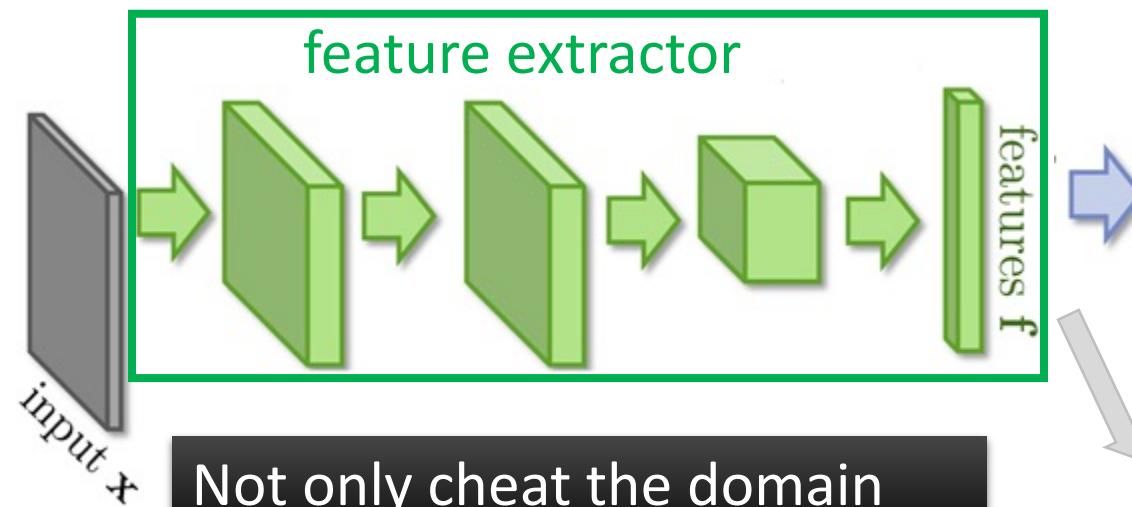
D tries to identify domains



Rq: Similar to GAN (coming soon)

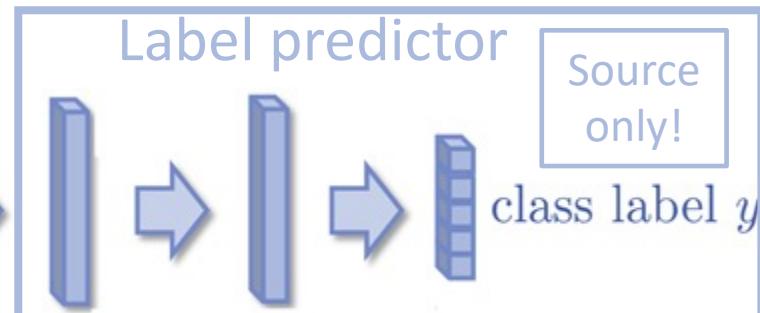
UDA strategy: 1/ domain-adversarial training  
2/ classification task (same for source  
and target here)

Maximize label classification accuracy +  
minimize domain classification accuracy



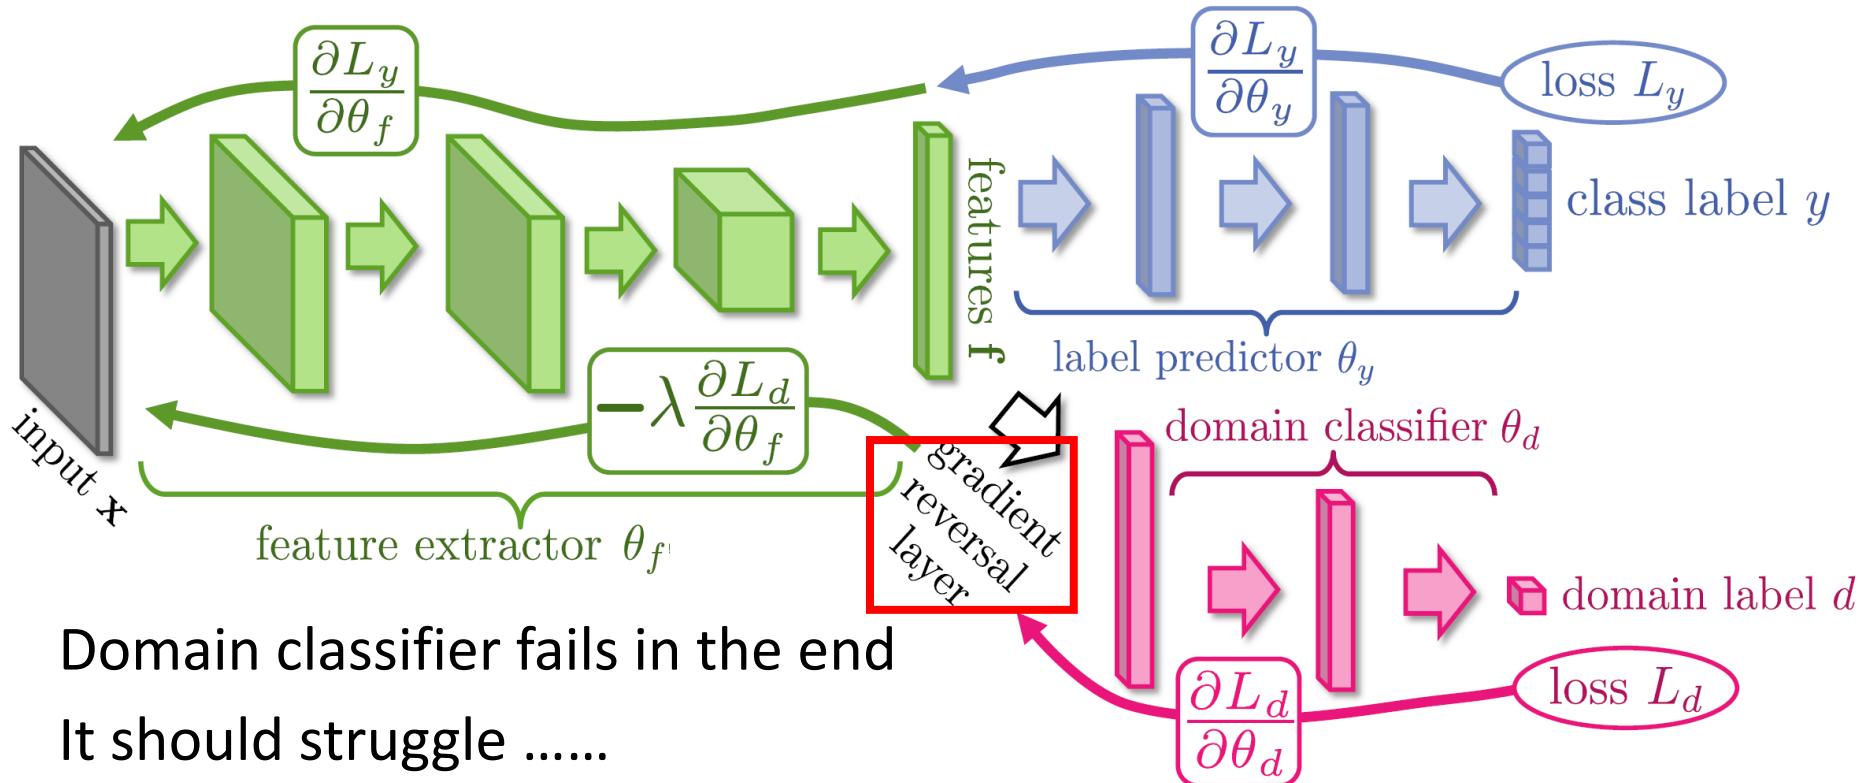
Not only cheat the domain  
classifier, but satisfying label  
classifier at the same time

Maximize label  
classification accuracy



Maximize domain  
classification accuracy

# UDA strategy: joint learning



Domain classifier fails in the end  
It should struggle .....

Yaroslav Ganin, Victor Lempitsky, Unsupervised Domain Adaptation by Backpropagation, ICML, 2015

Hana Ajakan, Pascal Germain, Hugo Larochelle, François Laviolette, Mario Marchand, Domain-Adversarial Training of Neural Networks, JMLR, 2016

# Domain-adversarial training



METHOD	SOURCE	MNIST	SYN NUMBERS	SVHN	SYN SIGNS
	TARGET	MNIST-M	SVHN	MNIST	GTSRB
SOURCE ONLY		.5749	.8665	.5919	.7400
SA (FERNANDO ET AL., 2013)		.6078 (7.9%)	.8672 (1.3%)	.6157 (5.9%)	.7635 (9.1%)
PROPOSED APPROACH		<b>.8149</b> (57.9%)	<b>.9048</b> (66.1%)	<b>.7107</b> (29.3%)	<b>.8866</b> (56.7%)
TRAIN ON TARGET		.9891	.9244	.9951	.9987

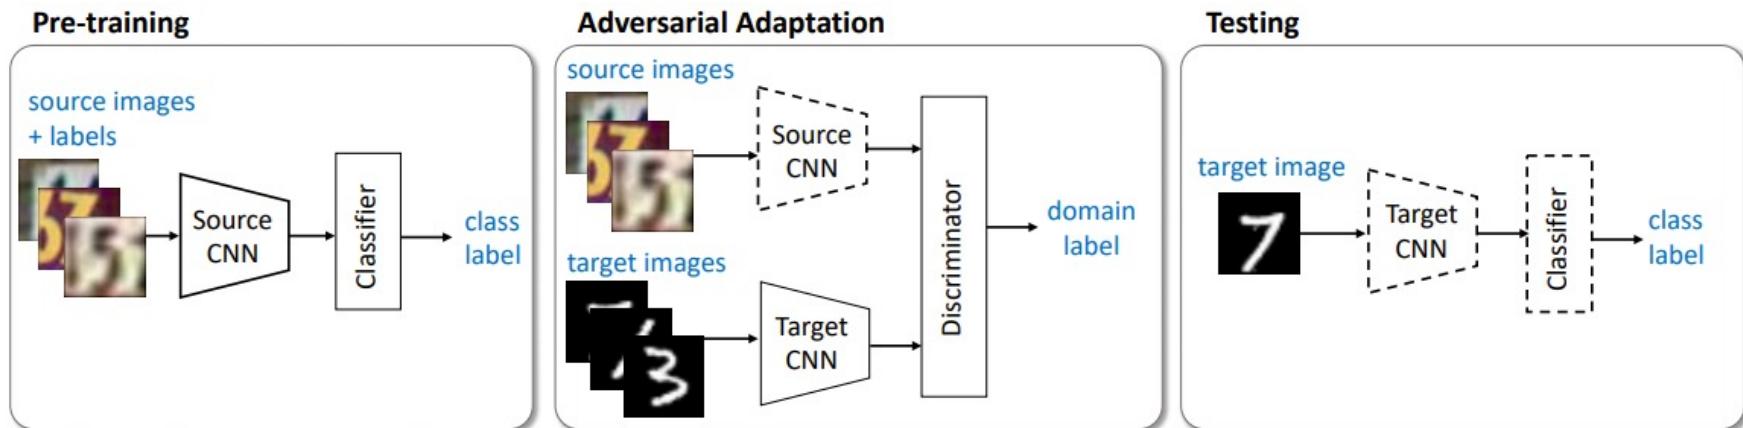
Yaroslav Ganin, Victor Lempitsky, Unsupervised Domain Adaptation by Backpropagation, ICML, 2015

Hana Ajakan, Pascal Germain, Hugo Larochelle, François Laviolette, Mario Marchand, Domain-Adversarial Training of Neural Networks, JMLR, 2016

# Domain adaptation

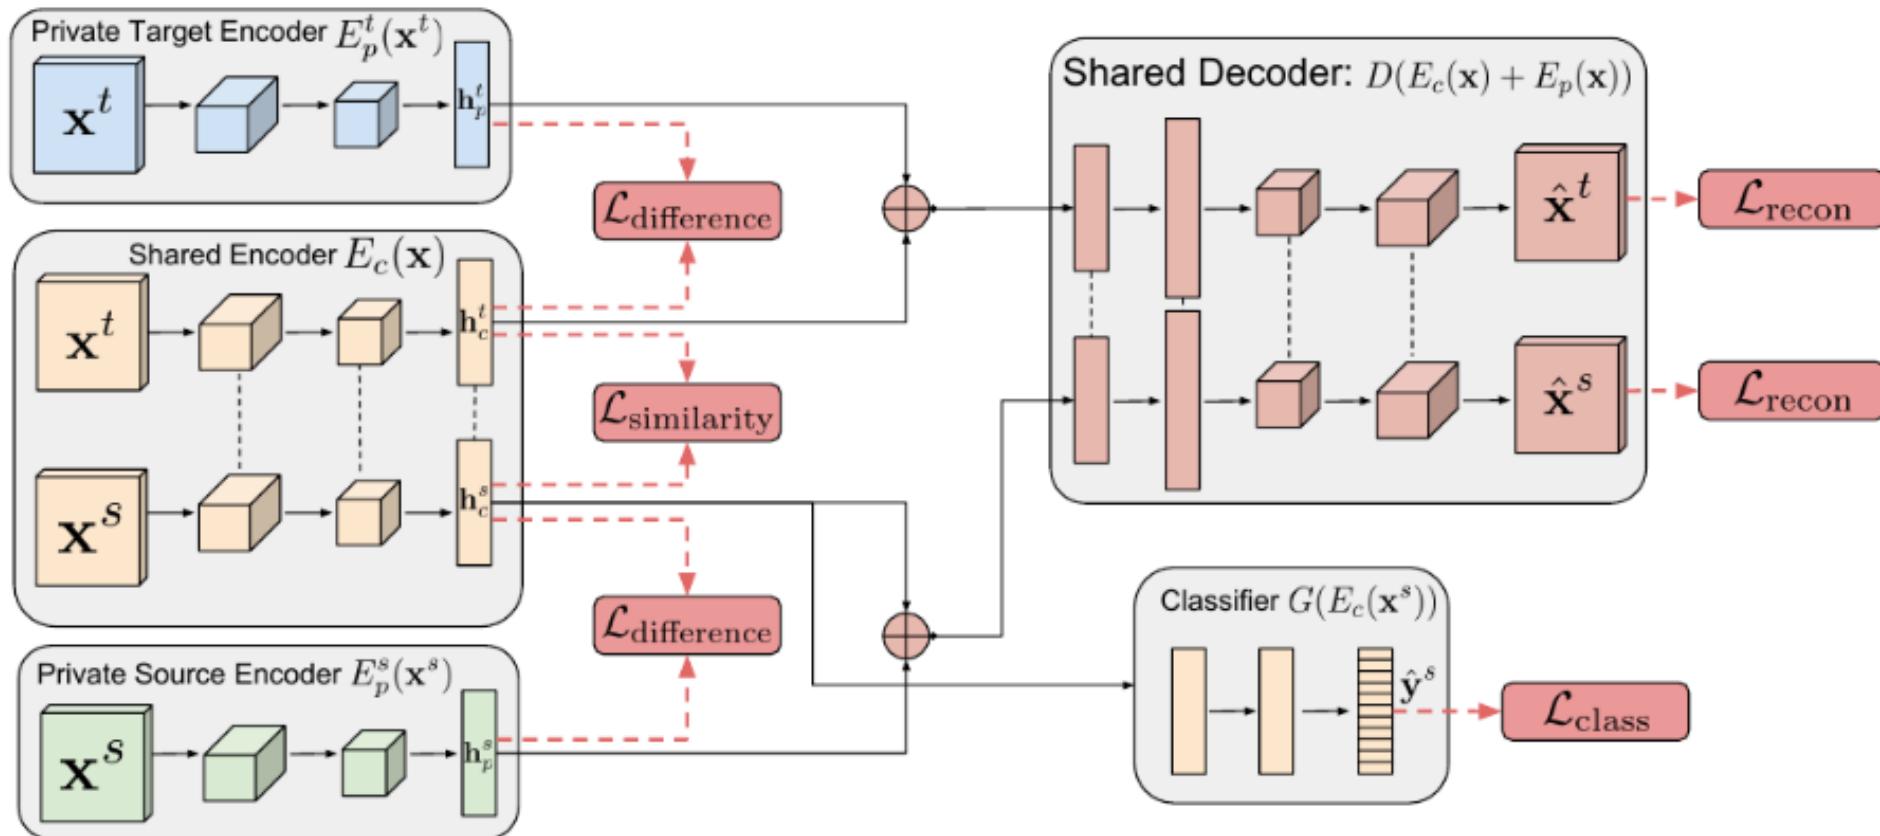
Main principle: diminish the domain shift in the learned features, encourage domain confusion

Another example: Adversarial Discriminative Domain Adaptation [Tzeng et al. 2017]



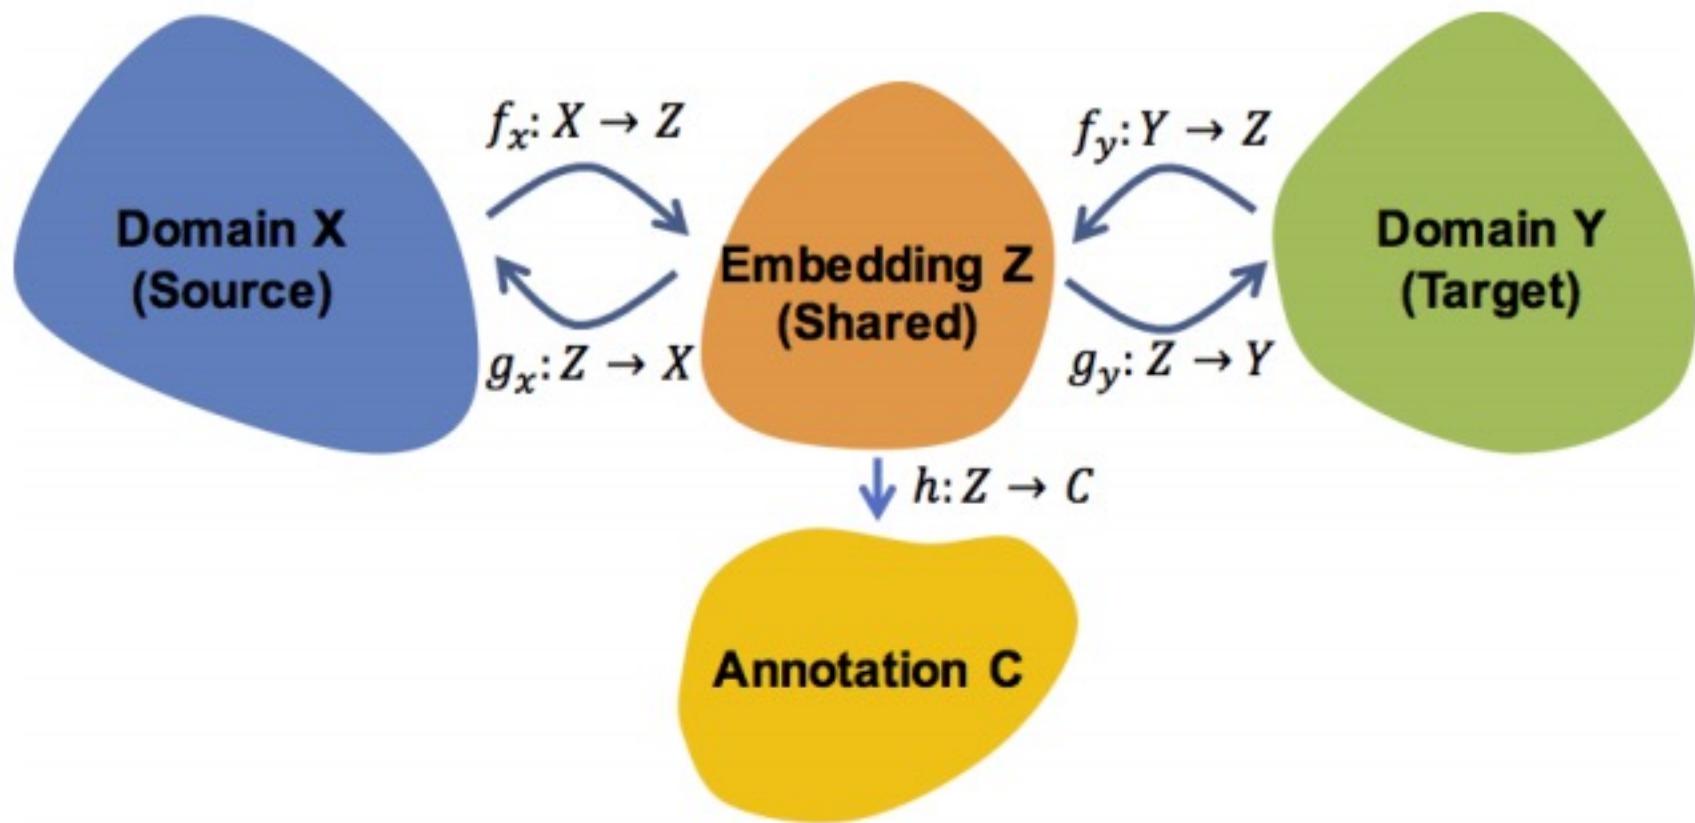
# Domain adaptation

## Other architecture



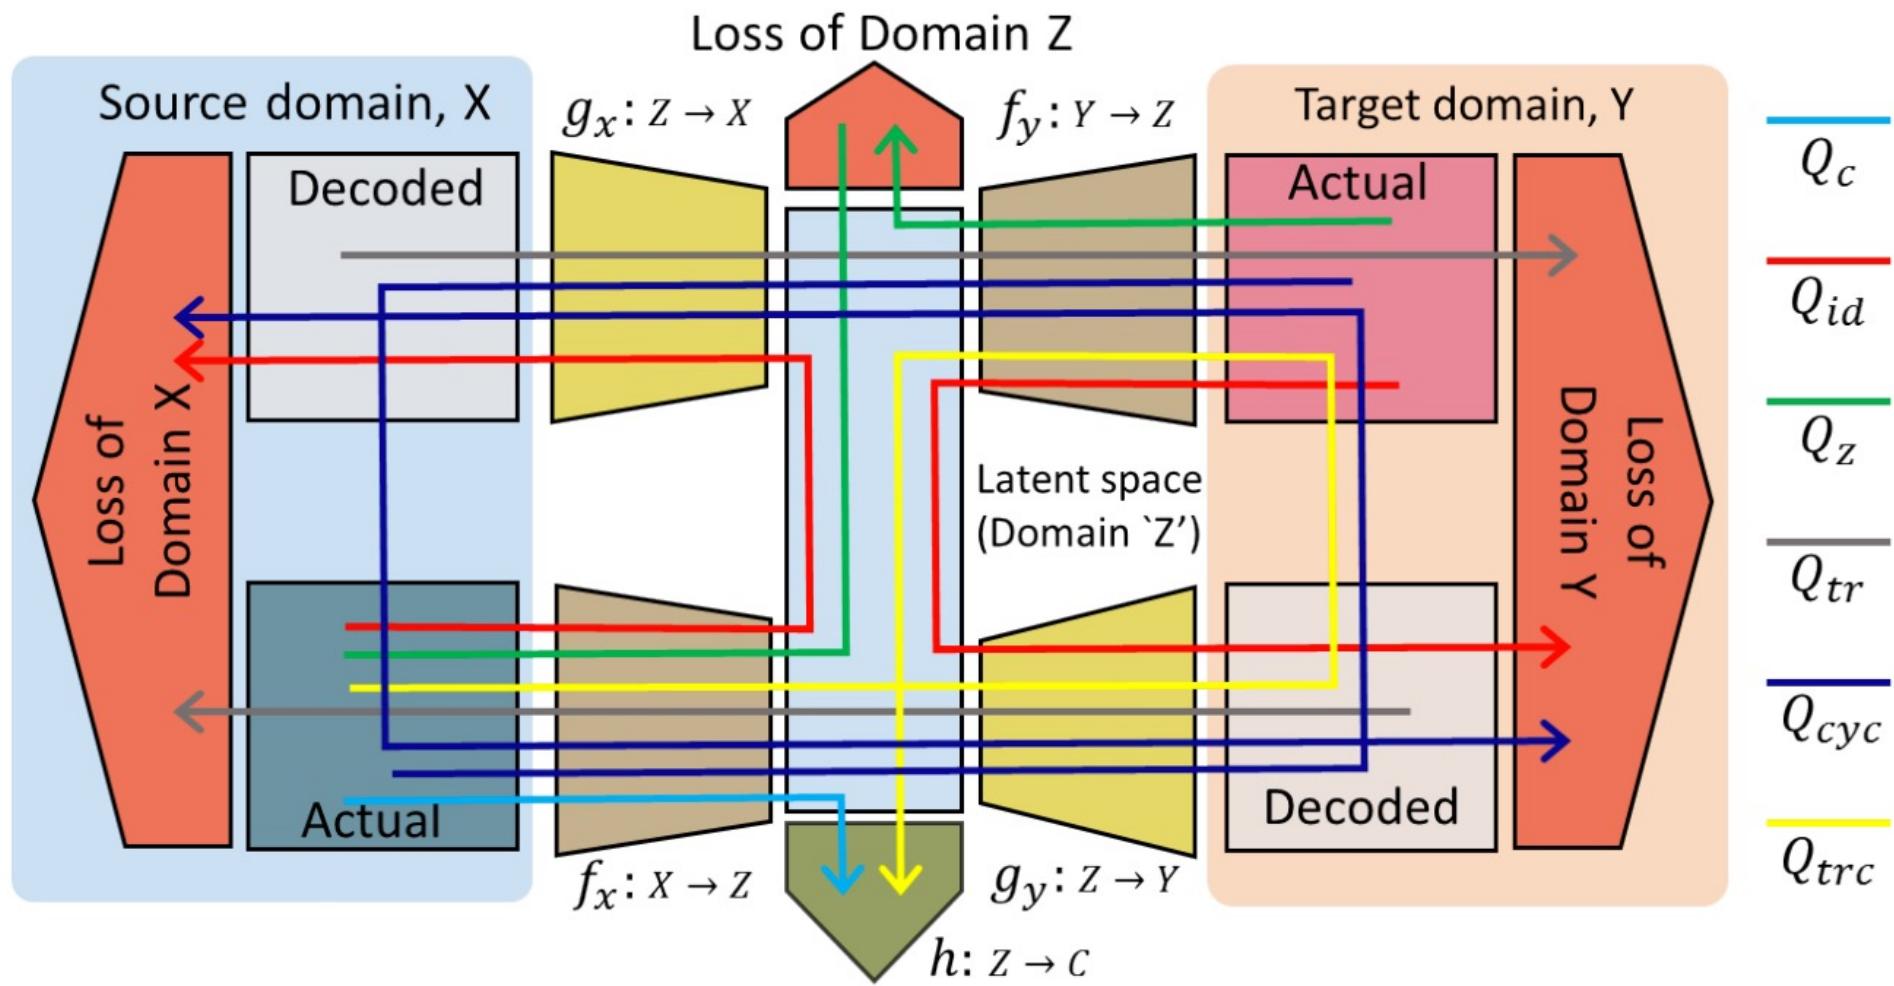
# Domain adaptation

Other architecture: Image translation for Domain adaptation [Murez 2017]



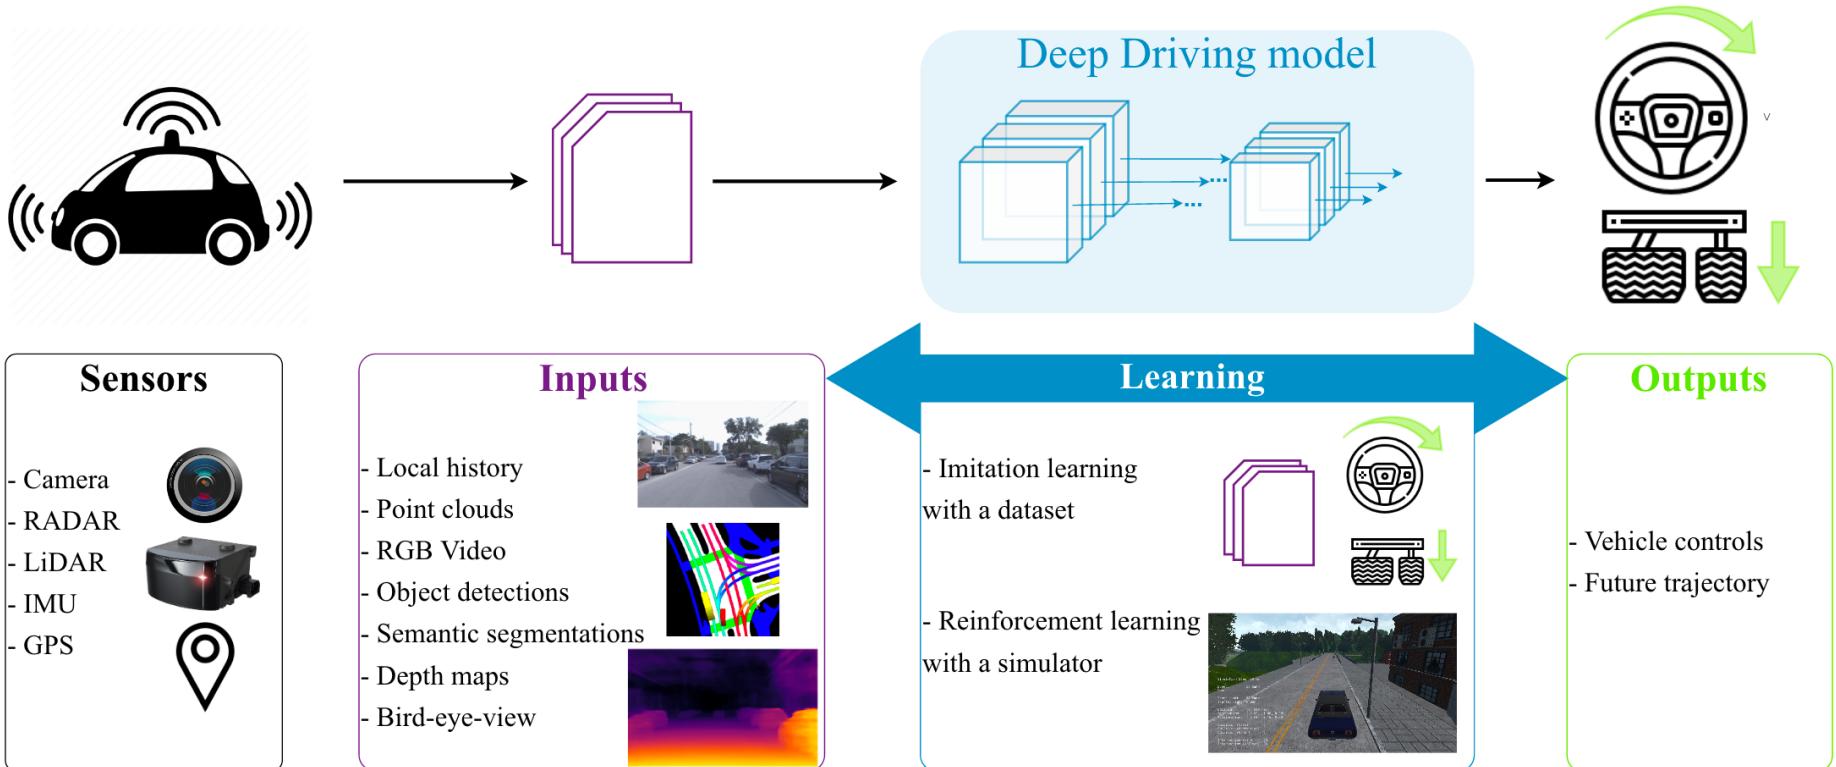
# Domain adaptation

Other architecture: Image translation for Domain adaptation [Murez 2017]



Use-Case: Domain adaptation for  
Autonomous driving

# Context: Neural network-based autonomous driving system framework



# Challenges for perception

## Multi-sensor perception

- Sensor fusion; Camera, radar and Lidar

## 3D dynamic understanding

- 3D object detection; Motion forecast; Intention prediction

## Frugal learning

- Training with limited data or supervision; Domain adaptation

## Reliability

- Robustness; Uncertainty estimation; Failure prediction

## Explainability

- Decision interpretation; Post-hoc or by-design

# Domain gap

Different, though *related* input data distributions

Source domain → Target domain



- Different weather, light, location, sensor's spec/setup

# Domain gap

Different, though *related* input data distributions

Source domain → Target domain



- Different weather, light, location, sensor's spec/setup

# Domain gap

Different, though *related* input data distributions

Source domain → Target domain



- Different weather, light, location, sensor's spec/setup

# Domain gap

Different, though *related* input data distributions

Source domain → Target domain



- Different weather, light, location, sensor's spec/setup

# Domain gap

Different, though *related* input data distributions

Source domain → Target domain

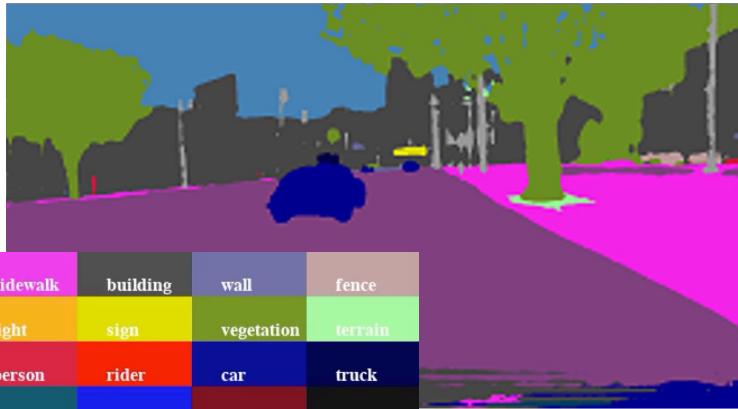


- Synthetic vs. real

# Domain gap for VISUAL SEGMENTATION

Different, though *related* input data distributions

Source domain → Target domain

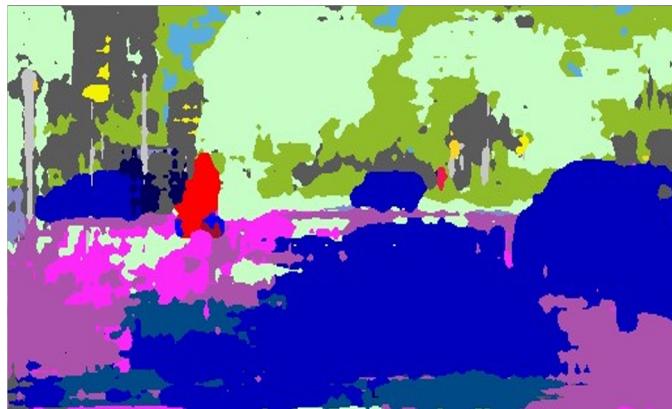


- Synthetic vs. real

# Domain gap

Different, though *related* input data distributions

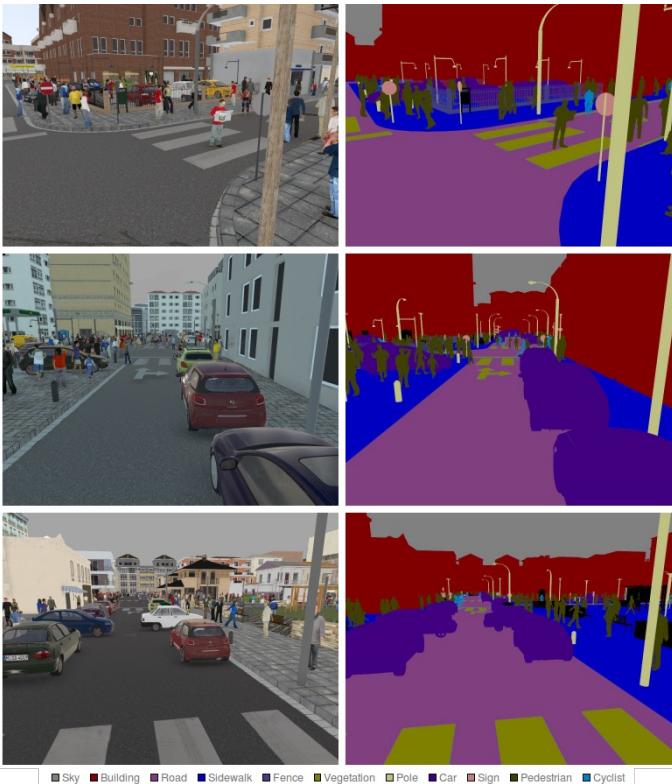
Source domain → Target domain



- Synthetic vs. real

# Unsupervised Domain Adaptation (UDA)

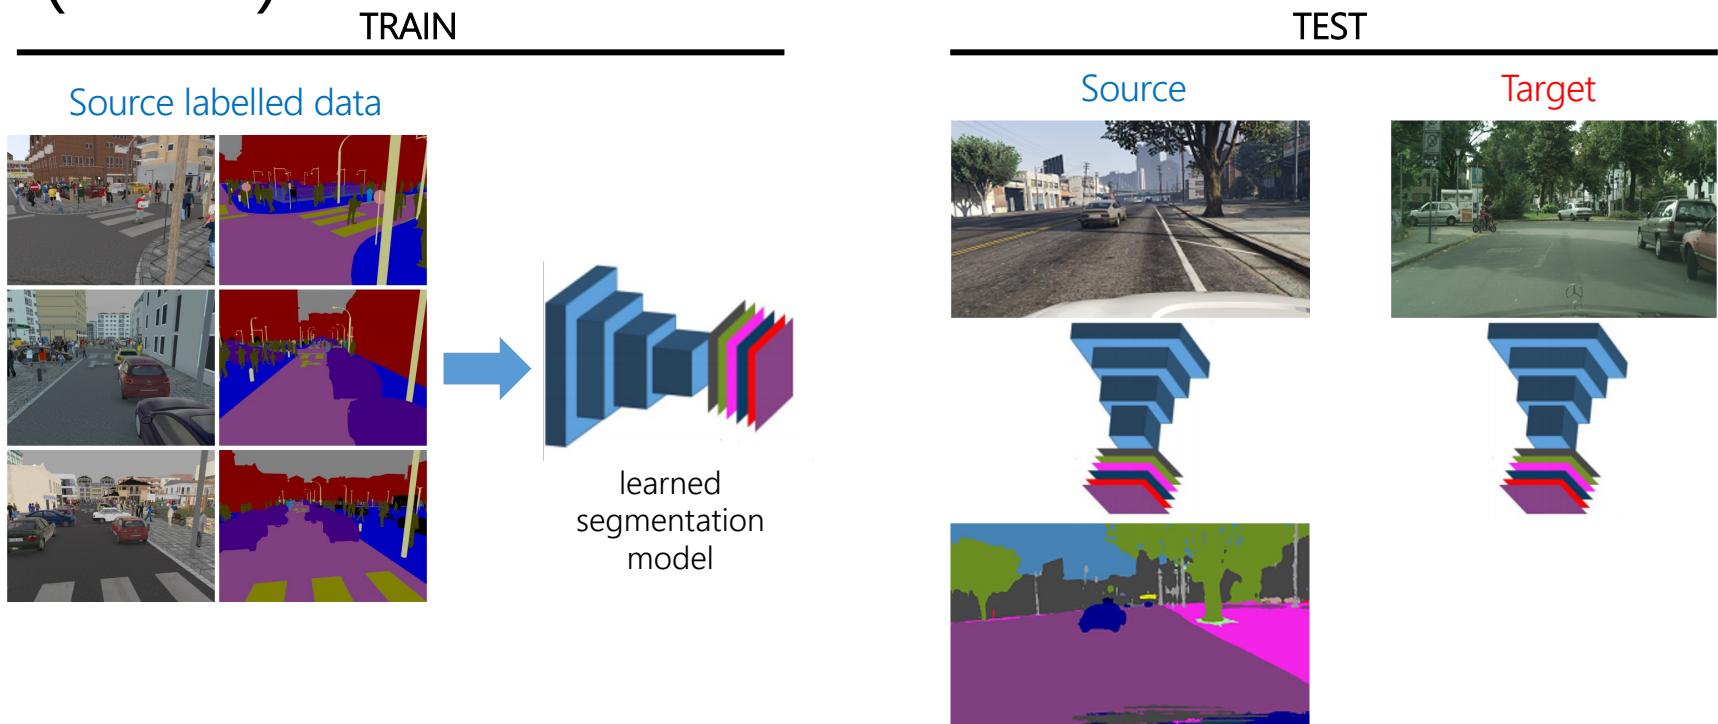
Labelled source domain data



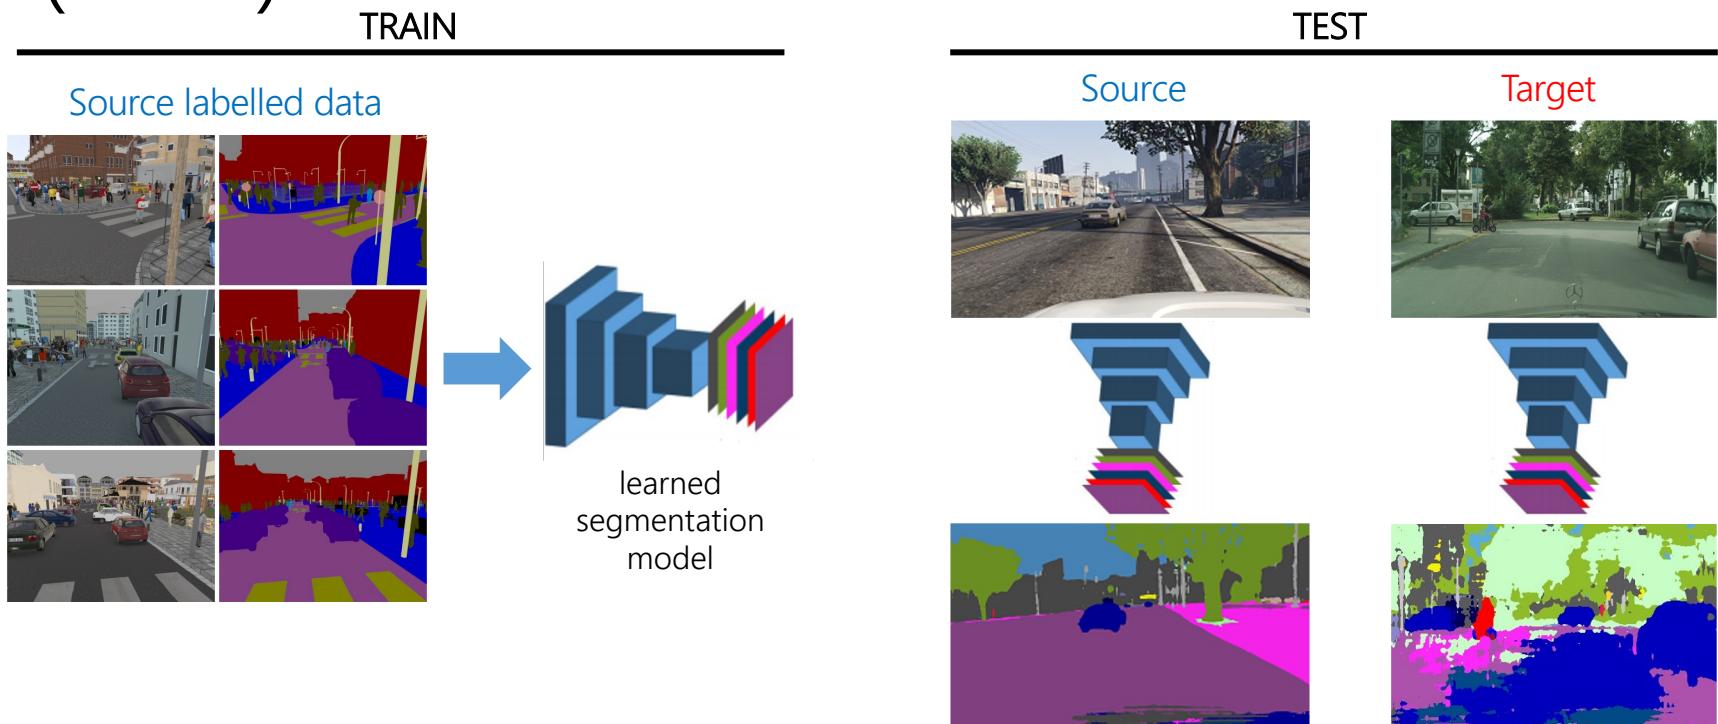
Unlabelled target domain data



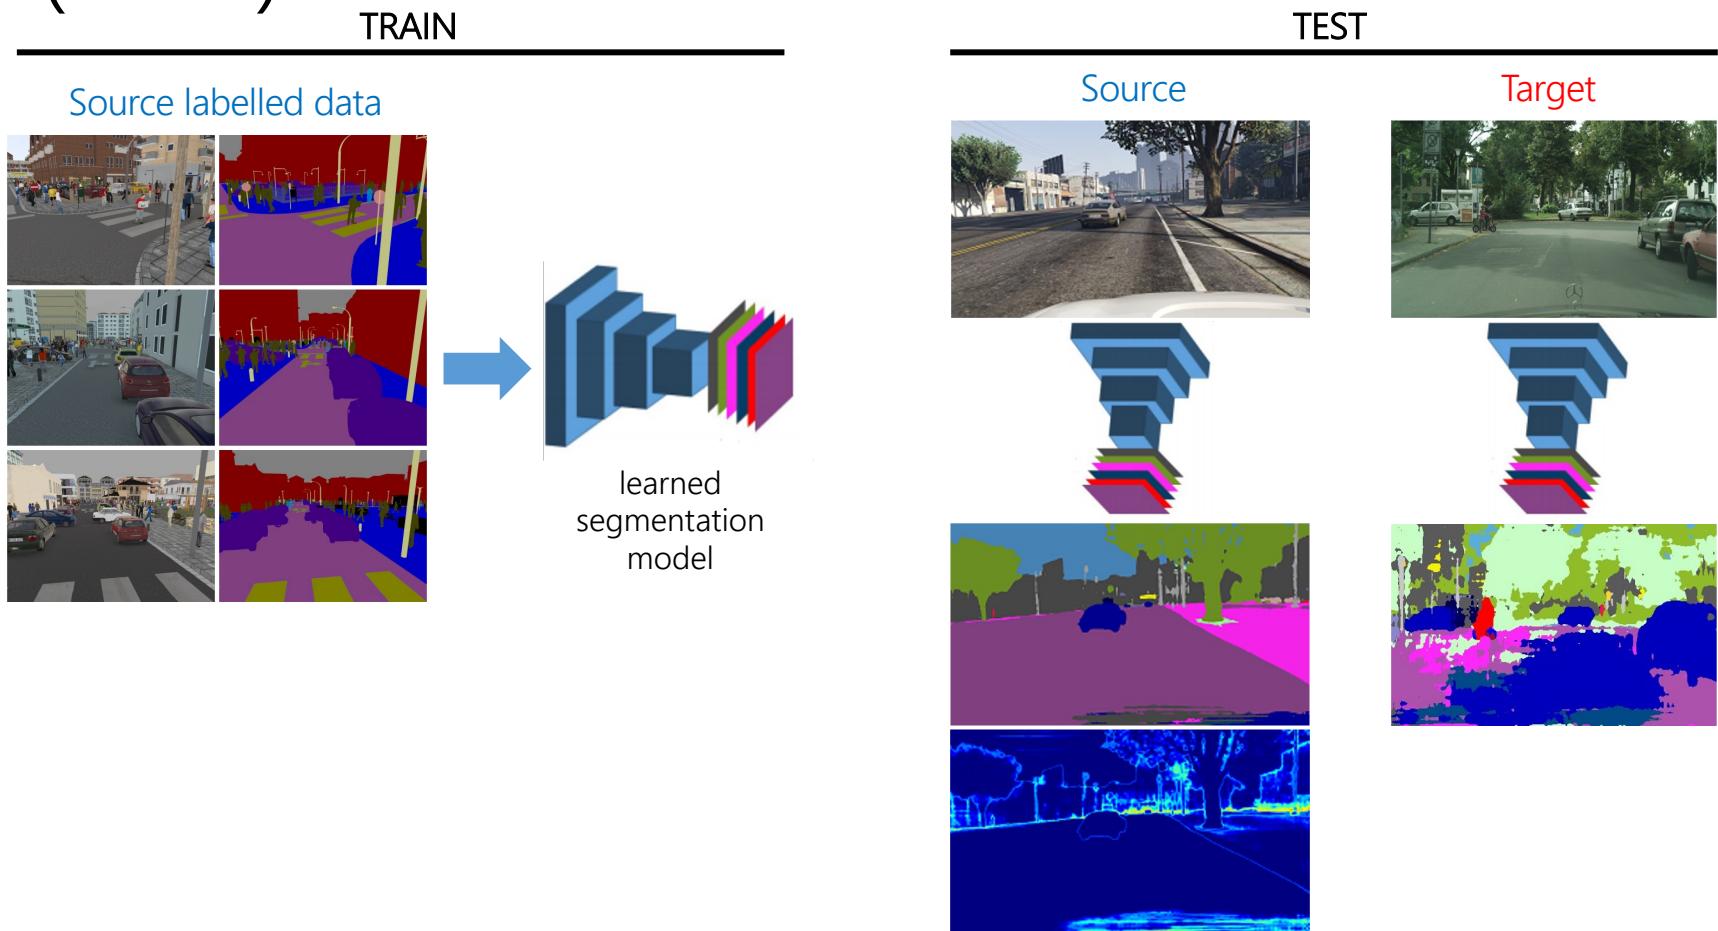
# Unsupervised Domain Adaptation (UDA)



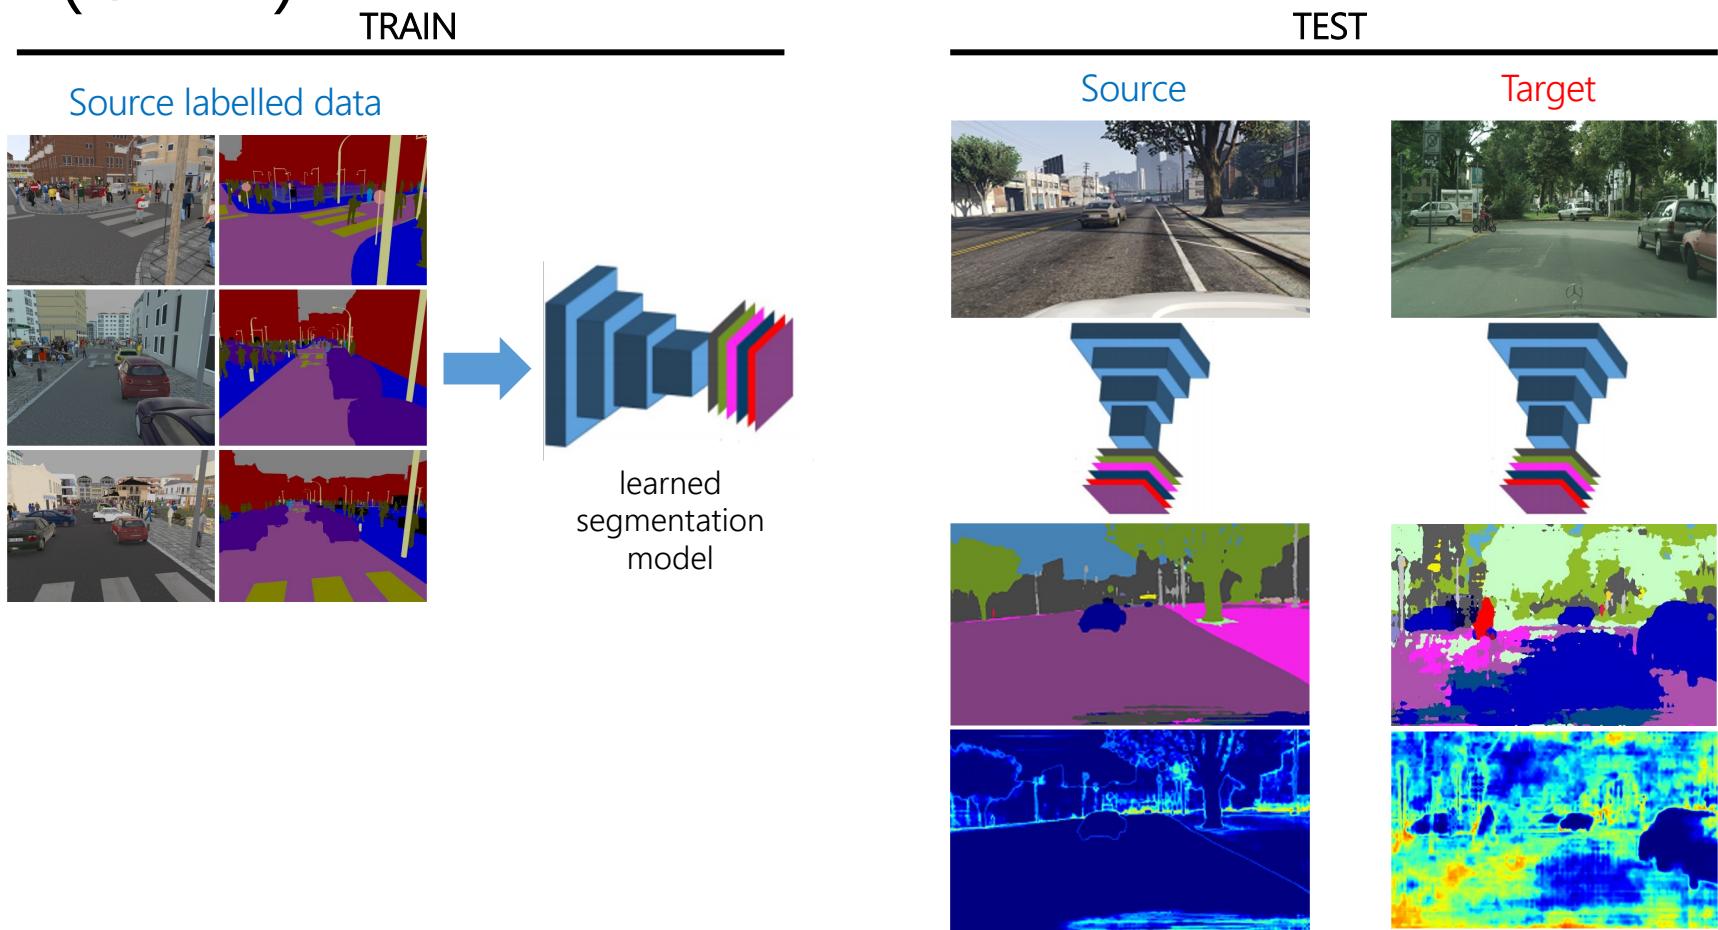
# Unsupervised Domain Adaptation (UDA)



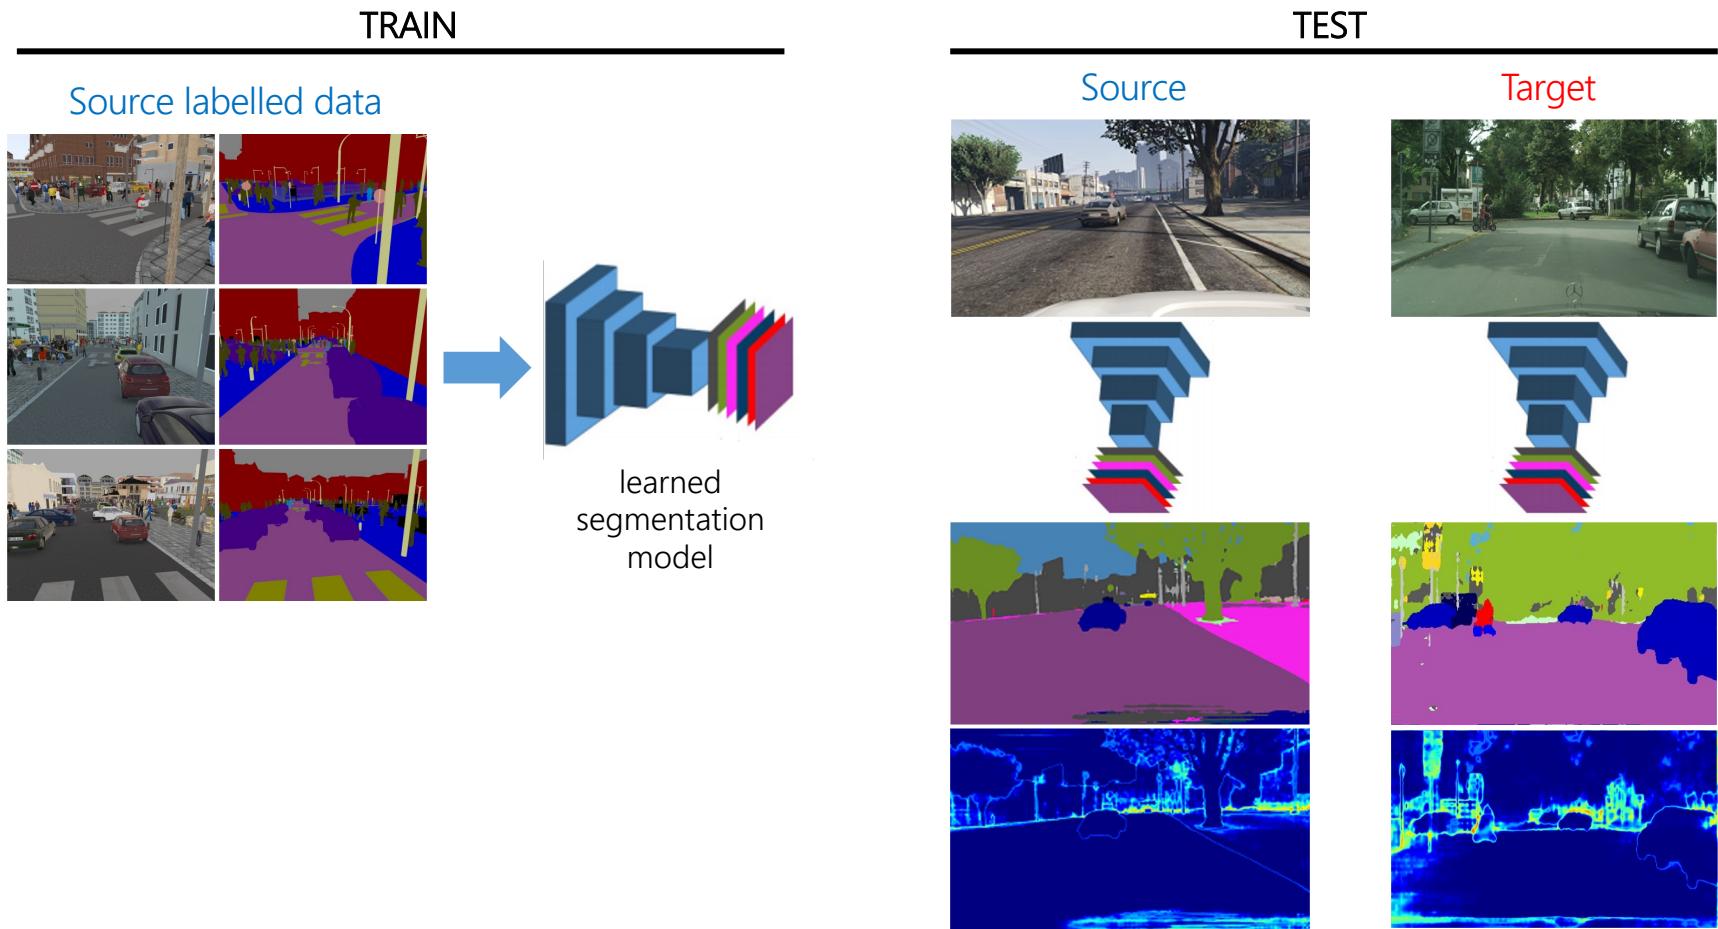
# Unsupervised Domain Adaptation (UDA)



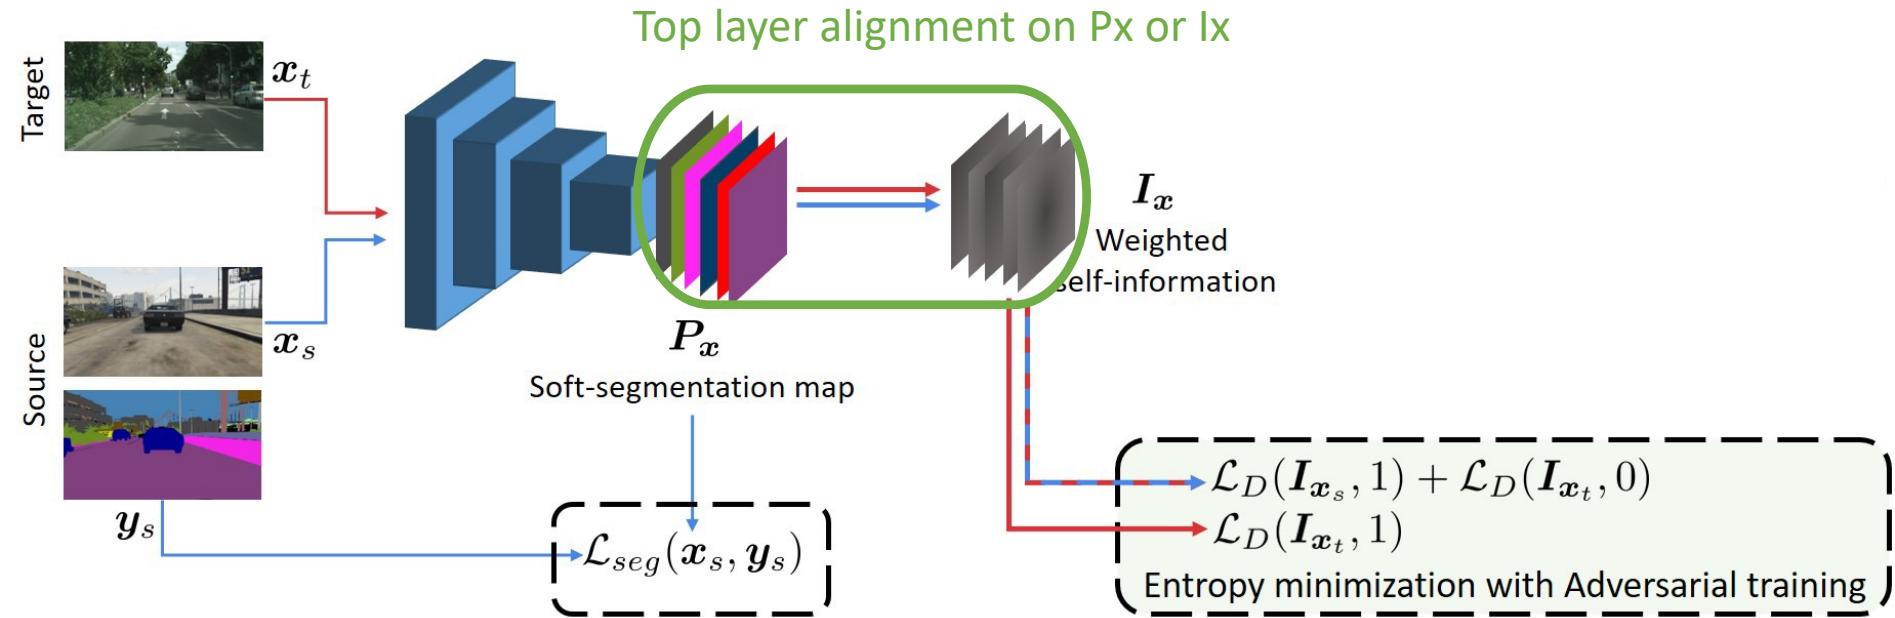
# Unsupervised Domain Adaptation (UDA)



# Expected results with UDA training



# Unsupervised Domain Adaptation (UDA)

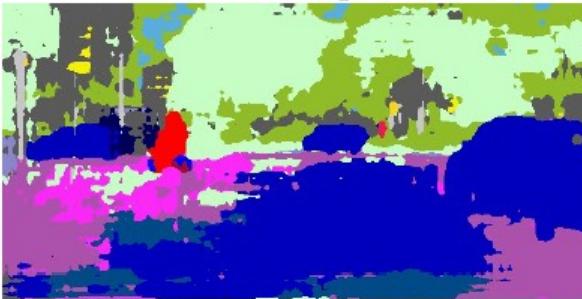


# Qualitative results

input image



without UDA

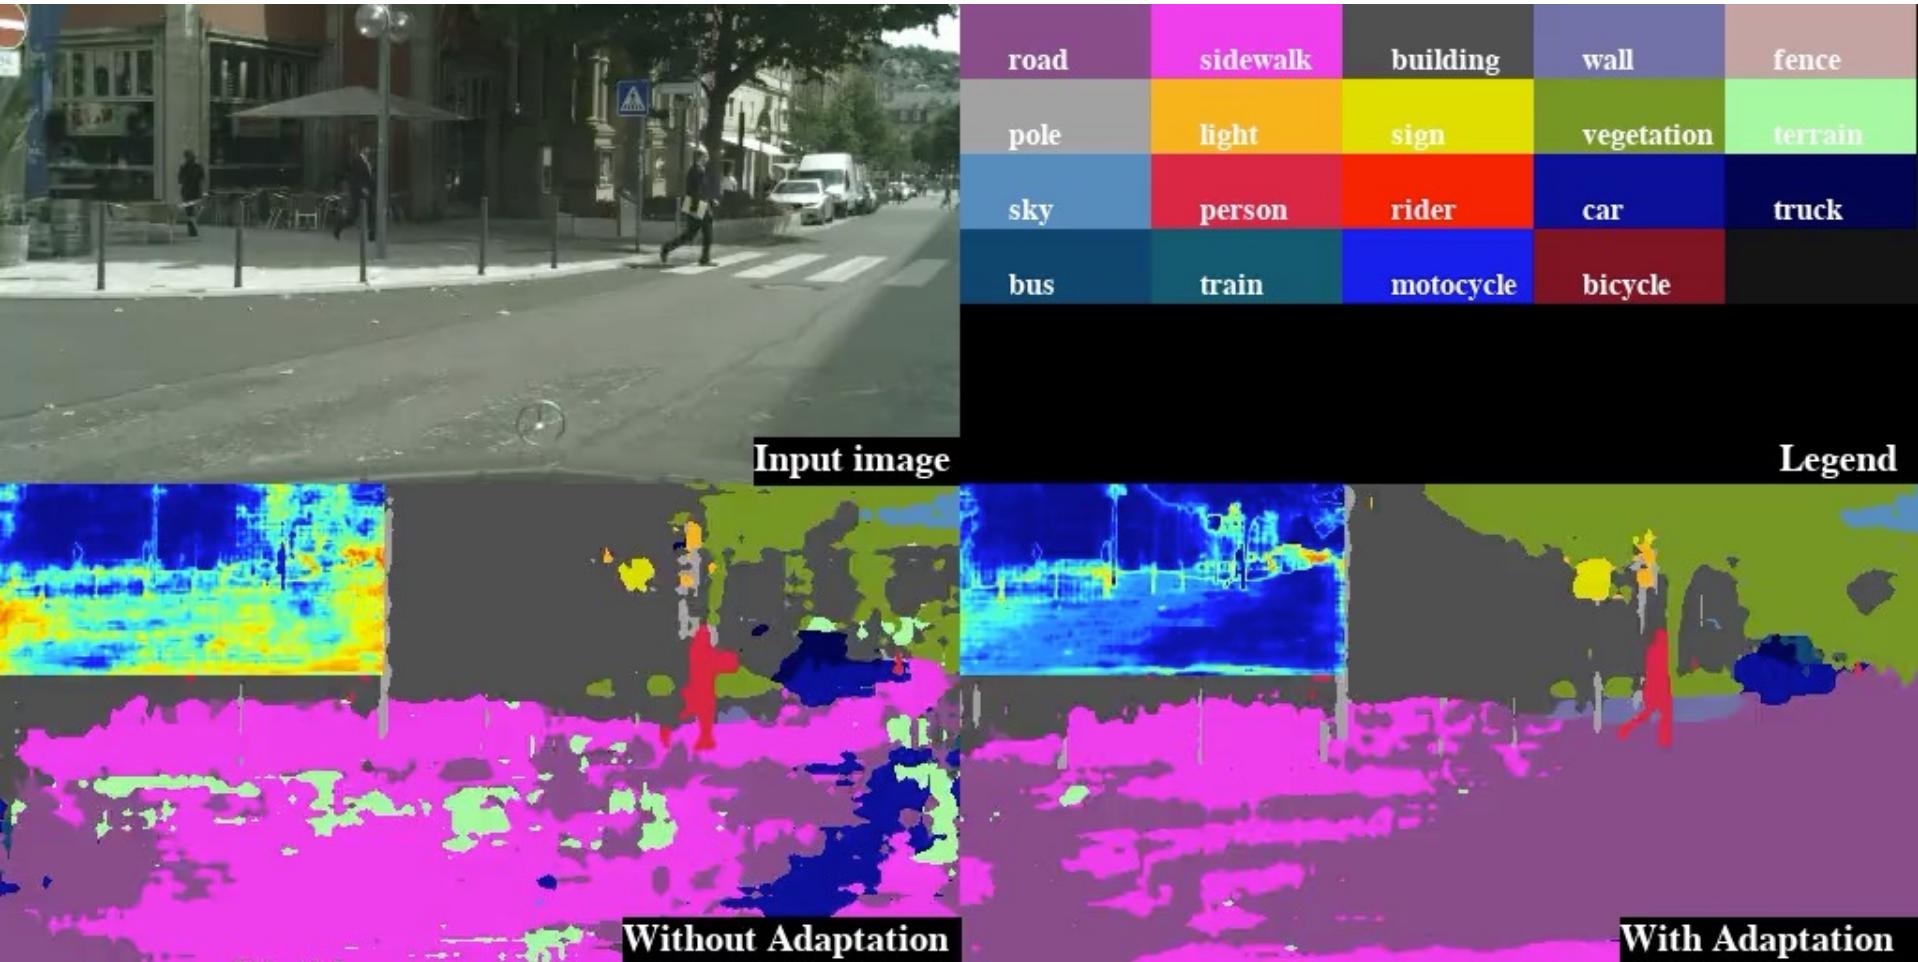


with UDA



road	sidewalk	building	wall	fence
pole	light	sign	vegetation	terrain
sky	person	rider	car	truck
bus	train	motorcycle	bicycle	

# UDA Results (with Adversarial Entropy)



# Extension: Zero shot + Domain adaptation



# Transfer Learning - Overview

		Source Data (not directly related to the task)	
		labelled	unlabeled
Target Data	labelled	Fine-tuning <i>Multitask Learning</i>	Not considered here
	unlabeled	Domain adaptation- adversarial training  <i>Zero-shot learning</i>	Not considered here

# Zero-shot Learning

- Source data:  $(x^s, y^s) \rightarrow$  Training data
- Target data:  $(\emptyset)$  usually same domain

Different tasks

*Training time :*  $x^s:$



$y^s:$  cat



...  
...

+ Class Information

*Test time  $x^t:$*

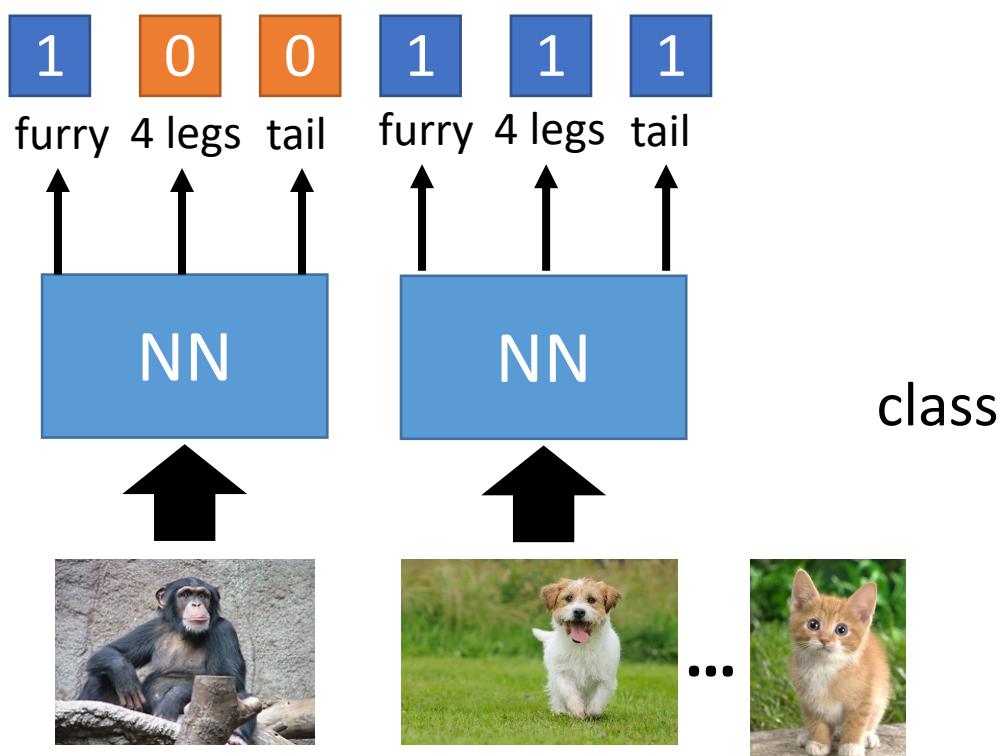


=> Fish class!

# Zero-shot Learning

- Representing each class by its attributes

## Training



+  
Database attributes

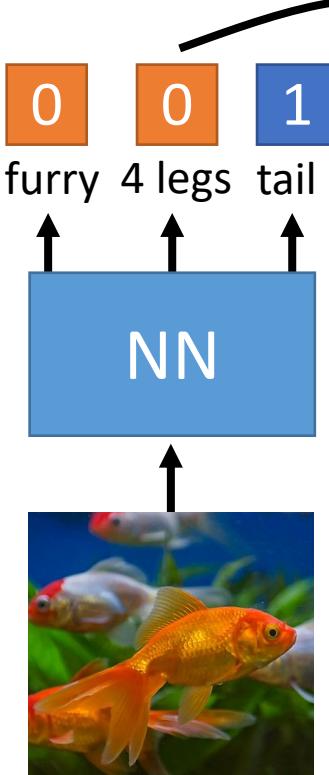
	furry	4 legs	tail	...
Dog	o	o	o	
Fish	x	x	o	
Chimp	o	x	x	
...				

sufficient attributes for one  
to one mapping

# Zero-shot Learning

- Representing each class by its attributes

## Testing



Find the class with the most similar attributes

	furry	4 legs	tail	...
Dog	o	o	o	
Fish	x	x	o	
Chimp	o	x	x	
...				

sufficient attributes for one to one mapping

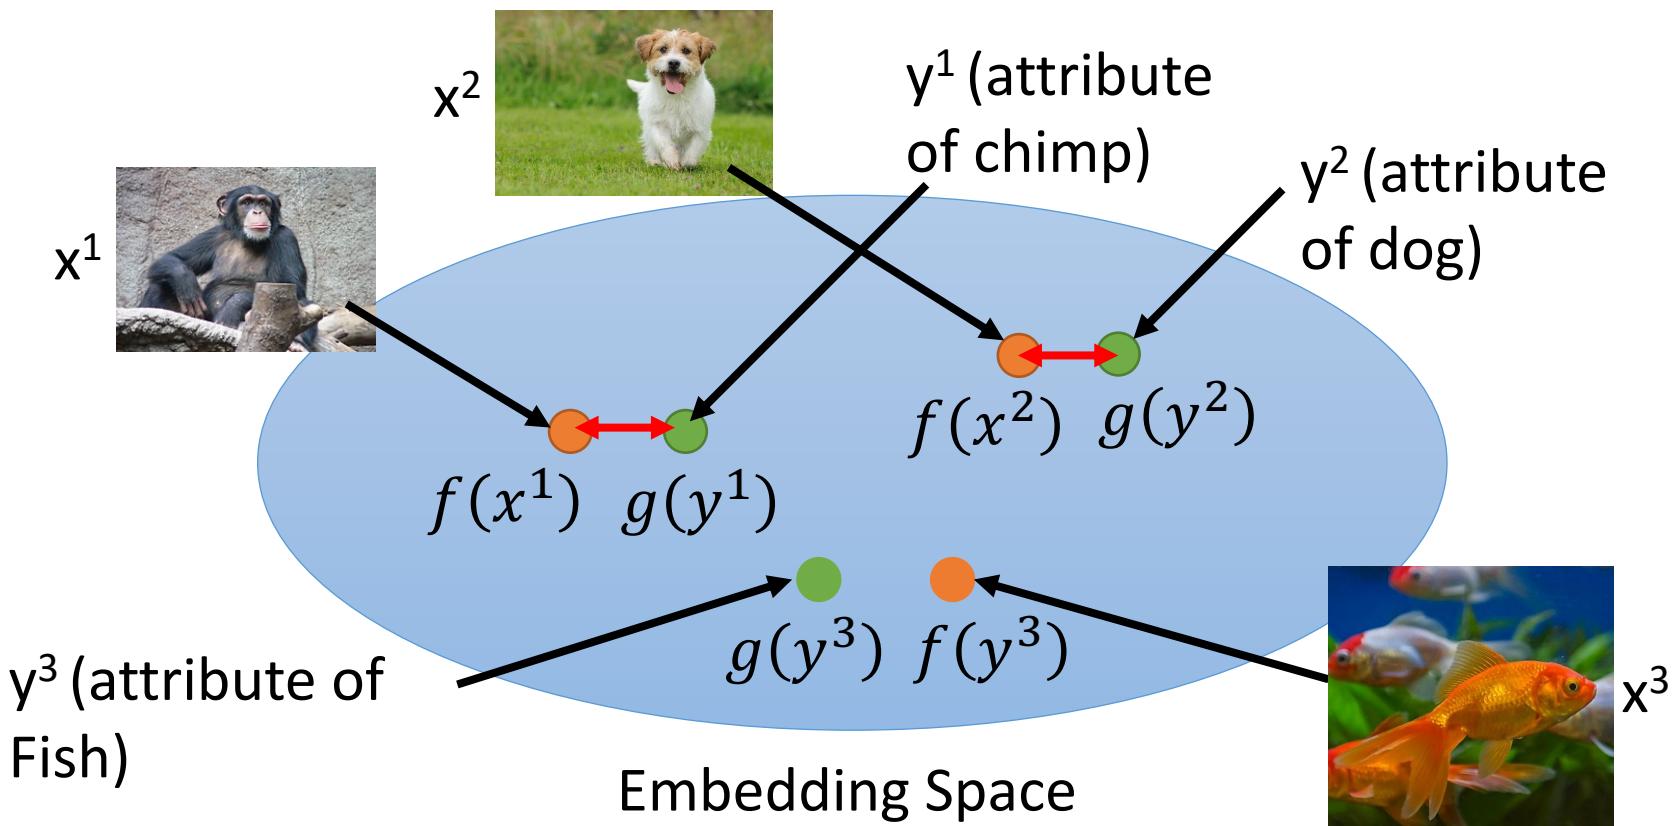
# Zero-shot Learning

What if we don't  
have attribute  
database

- Attribute embedding + class (word name) embedding

# Zero-shot Learning

- Attribute embedding



$f(\cdot)$  and  $g(\cdot)$  can be NN.  
Training target:

$f(x^n)$  and  $g(y^n)$  as close as possible

$y^i$  are linked together by a class relationship (e.g. class name embedding as W2v)